

# Assignment 4 & Benchmarks

Introduction to Database Systems

DataLab

CS, NTHU

# Outline

- Assignment 4
- Benchmarks
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- Example Results
- Benchmarking with Scripts

# Outline

- Assignment 4
- Benchmarks
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- Example Results
- Benchmarking with Scripts

# Outline

- Assignment 4
- **Benchmarks**
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- Example Results
- Benchmarking with Scripts

# The Micro-Benchmark

- Two types of transactions.
  - **Read-only** transaction => reads 10 records.
  - **Read-write** transaction => reads and updates 10 records.
  - The ratio is controlled by **RW\_TX\_RATE**.
- The data set is split into two parts.



- 1 is chosen from hot set, 9 are chosen from cold set.
- The number of hot records is control by **HOT\_CONFLICT\_RATE**.

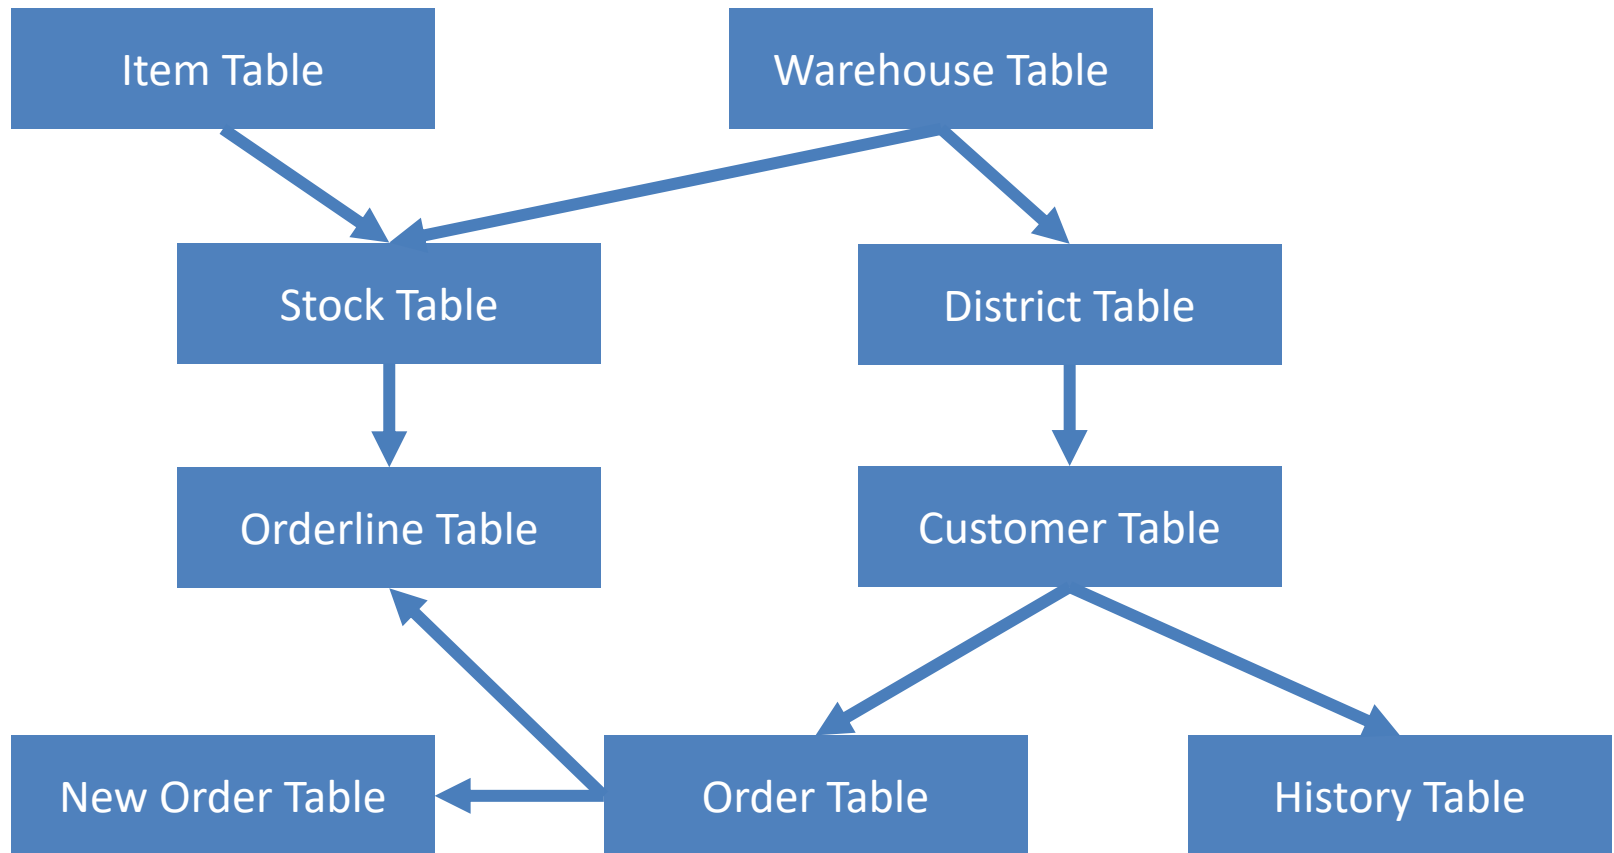
# Outline

- Assignment 4
- **Benchmarks**
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- Example Results
- Benchmarking with Scripts

# The TPC-C Benchmark

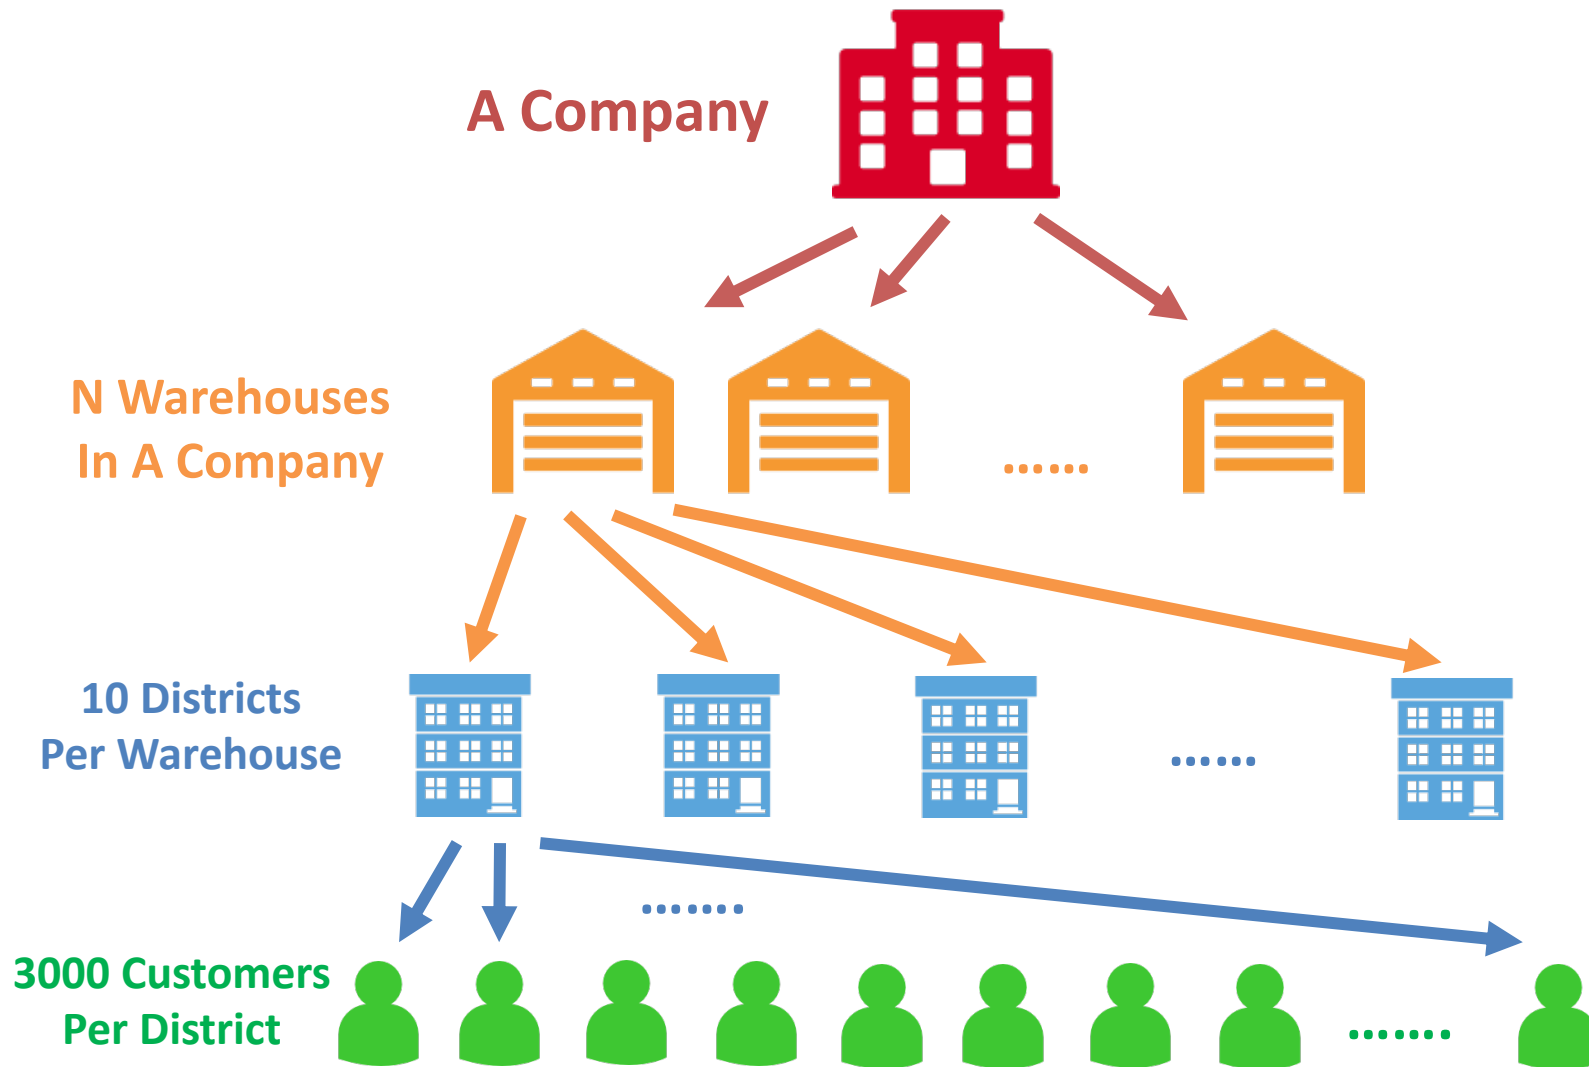
- The TPC-C benchmark is a industry-standard benchmark purposed by TPC (Transaction Processing Council).
  - There are also TPC-A, TPC-B, TPC-E, TPC-H.
- It simulates a warehouse management system.
  - Tree-structured: almost all records are related to a warehouse record.
  - Easy-to-partition: good for a distributed DBMS.

# Database Architecture





# Warehouses (Tree-Structured)



# Orders

A Customer



A New Order



An Order



An Orderline

An Orderline

An Orderline

An Orderline

# Types of Transactions

- New Order
  - 23 reads, 11 updates, 12 inserts in average.
- Payment
  - 4 reads, 3 updates, 1 insert.
- Stock Level
- Order Status
- Delivery

# Outline

- Assignment 4
- Benchmarks
  - The Micro-benchmark
  - The TPC-C Benchmark
- **Guidelines for Experiments**
- Example Results
- Benchmarking with Scripts

# Guidelines for Experiments

- Think about what settings can highlight your improvement.
- Make sure there is no other CPU-intensive programs running on the testing machines.
- Put the server and the client on different machine if you can.
- Use stored procedures.
- Using a fresh database every time.
- Find best # of RTEs before real experiments.
  - Which give you highest throughput.
- Throughput is a more important indicator for concurrency than latency.
- Draw you results as line plots or histograms in the report.

# Outline

- Assignment 4
- Benchmarks
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- **Example Results**
- Benchmarking with Scripts

# Example Results for the Micro-benchmarks

- Settings
  - RTE = 10
  - RW Tx Rate = 0.5
  - Conflict Rate = 0.001
- Throughputs (txs/min)

Buffer Size	Basic Version	Optimized Version	Speed Up
100000	111558	174521	56%
100	39285	75164	91%

# Outline

- Assignment 4
- Benchmarks
  - The Micro-benchmark
  - The TPC-C Benchmark
- Guidelines for Experiments
- Example Results
- **Benchmarking with Scripts**



# Why Do We Need Scripts?

1. To setup the system quickly.
2. To deploy and benchmark the system in different machines.
3. The environment may not have Eclipse!

# Check Your Environment

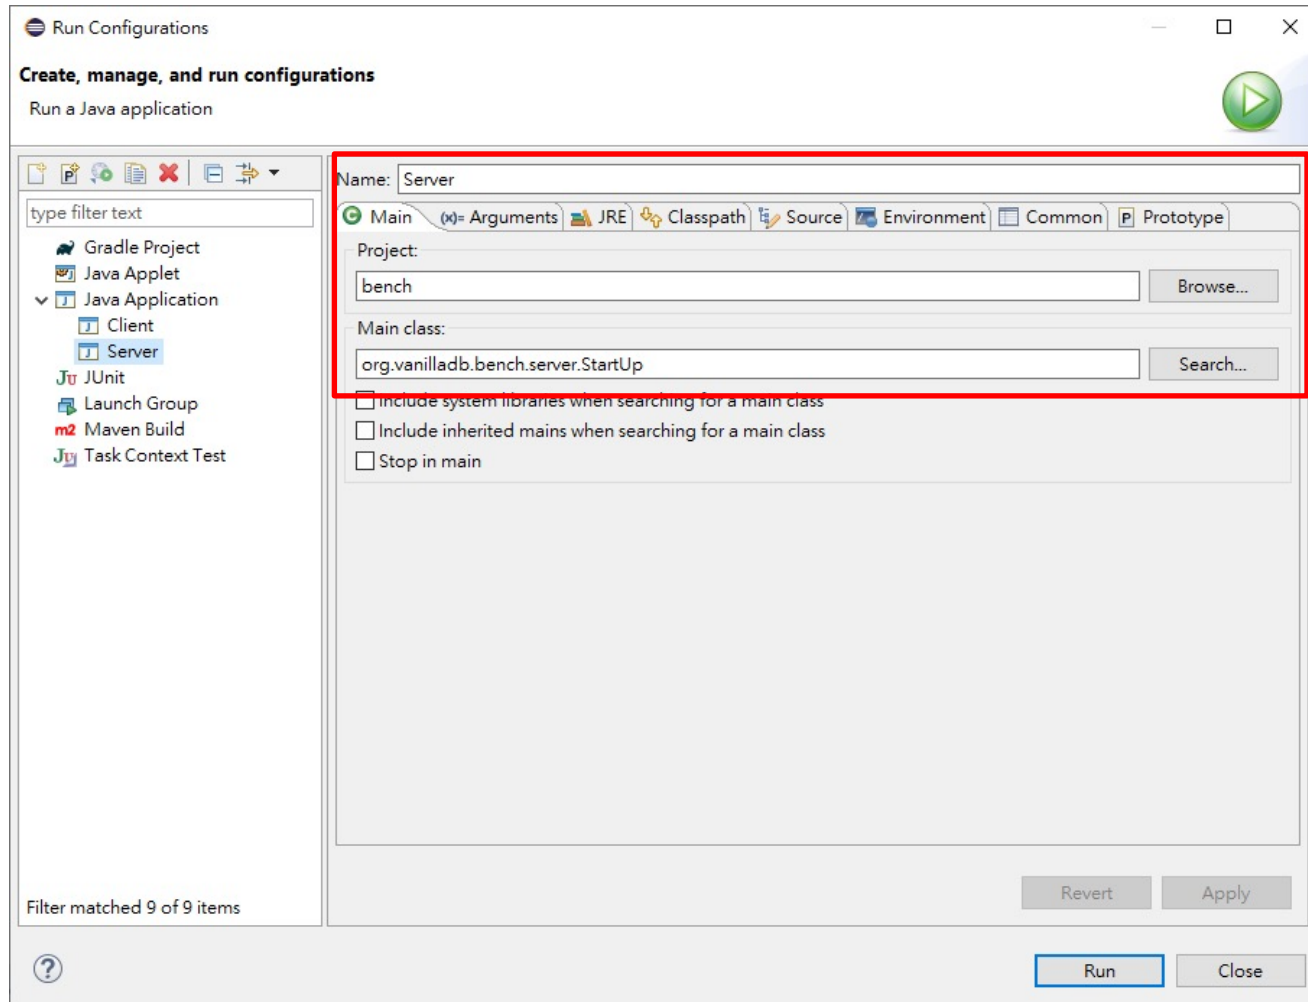
- Requirements
  - Bash
    - Which you may have had if you are using Unix, Unix-like systems or have installed Git on Windows.
  - Java in your system path

```
> java -version
```

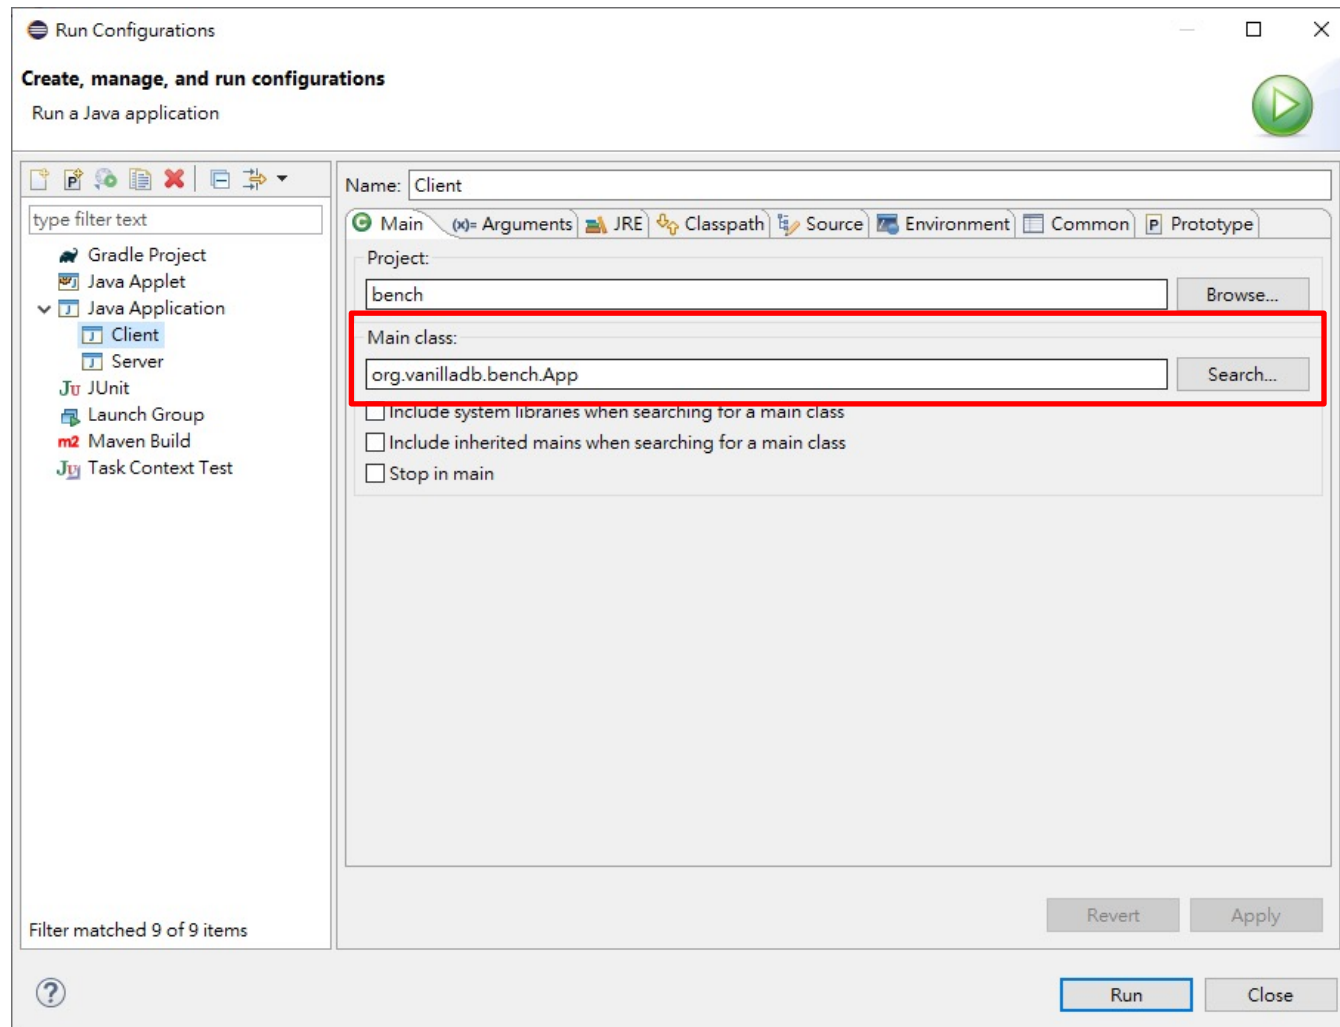
# Package Your Code

- We use Eclipse built-in tools.
- Steps
  1. Setup run configurations for jars.
  2. Export the project.

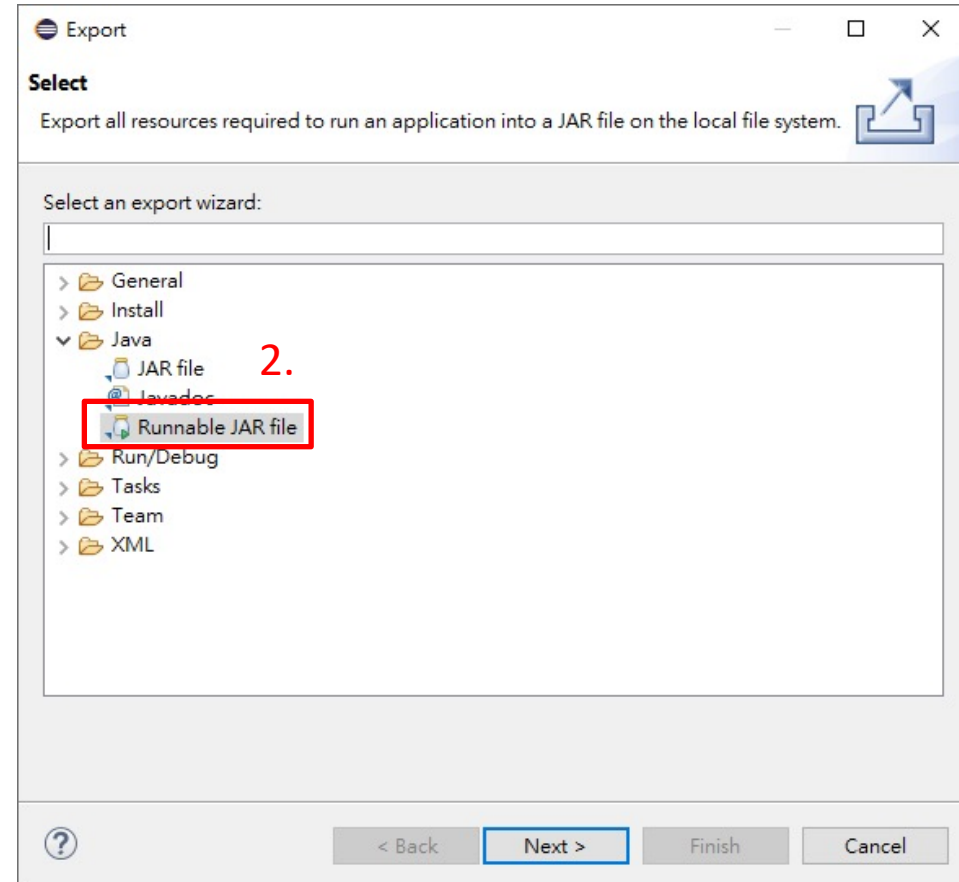
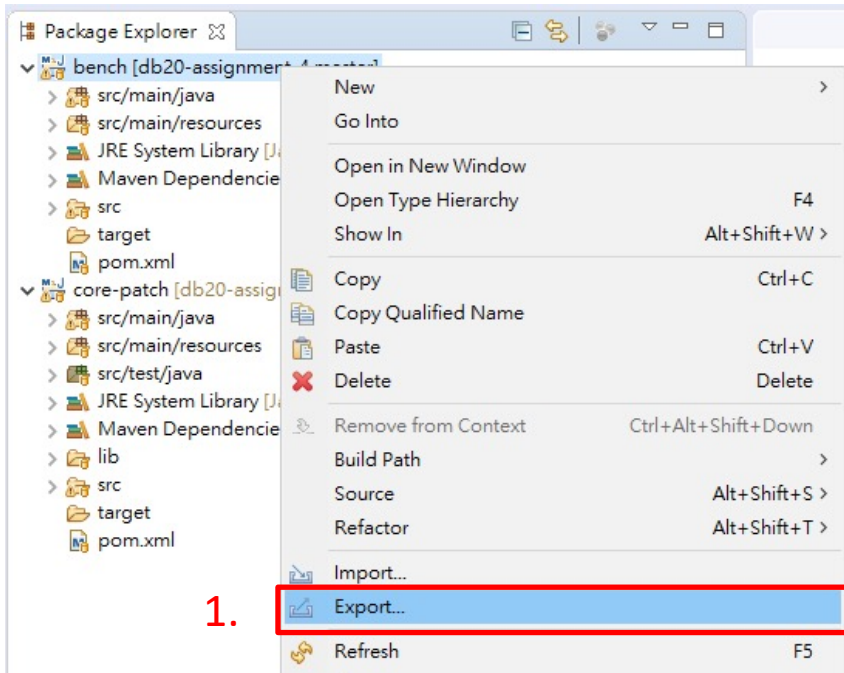
# Setup Run Configurations - Server



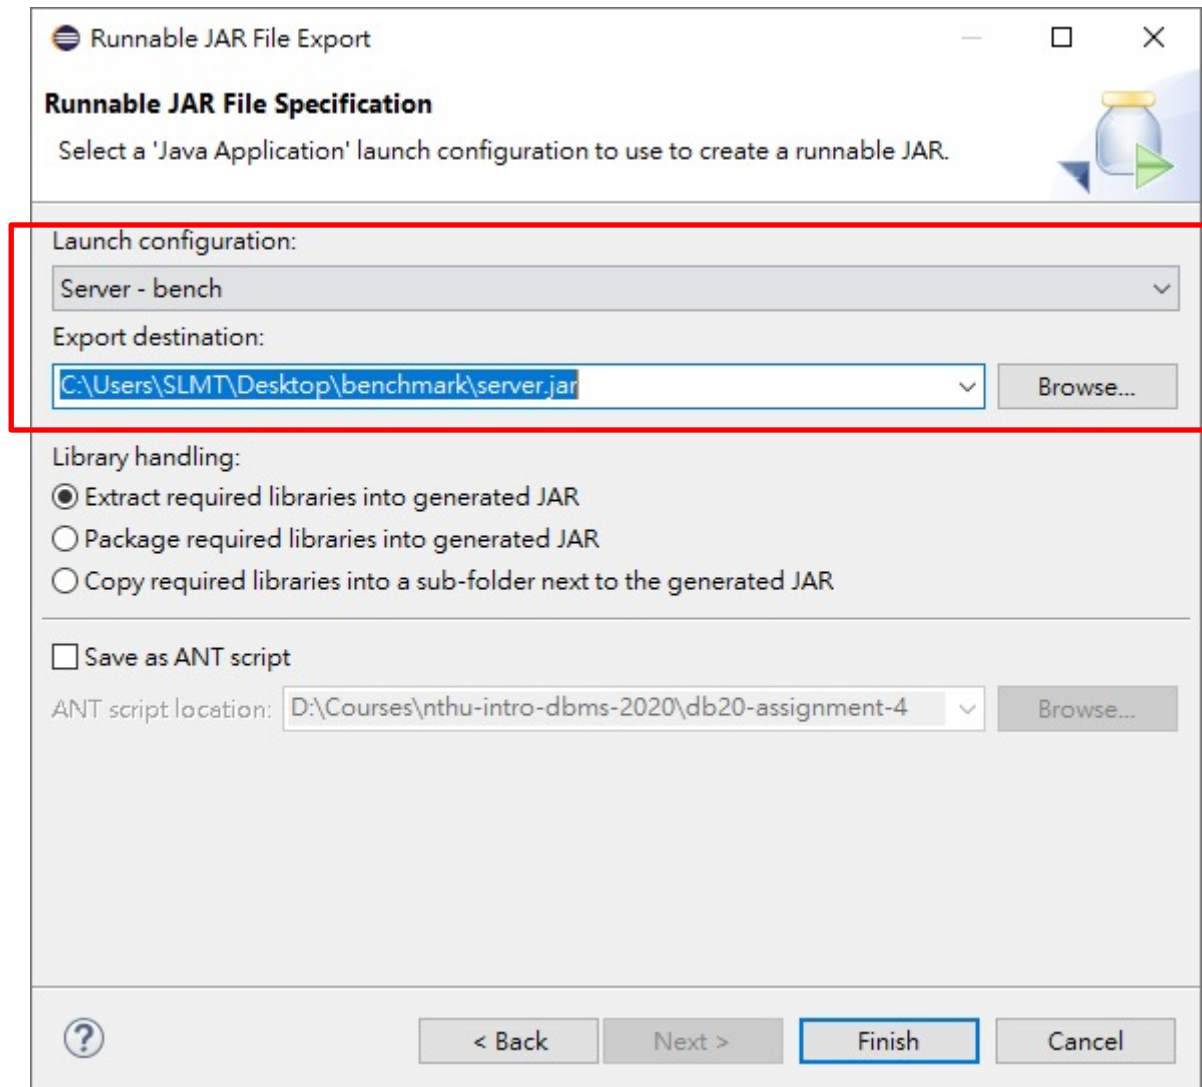
# Setup Run Configurations - Client



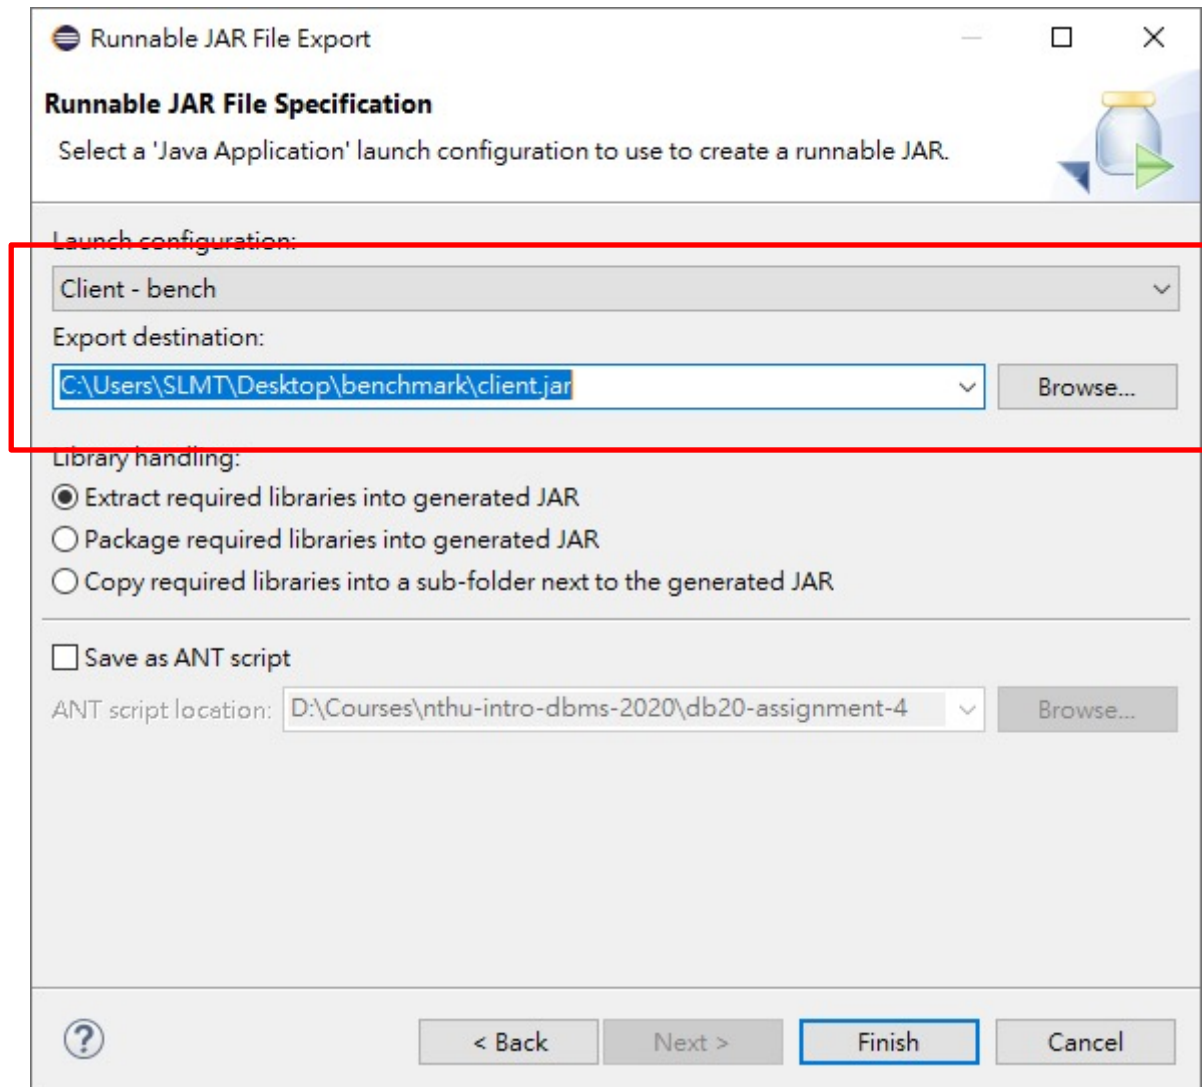
# Export Runnable Jars



# Export Runnable Jars - Server



# Export Runnable Jars - Client



The image shows a 'Runnable JAR File Export' dialog box. A red rectangle highlights the 'Launch configuration' dropdown (set to 'Client - bench') and the 'Export destination' text field (containing 'C:\Users\SLMT\Desktop\benchmark\client.jar').

**Runnable JAR File Export**

Select a 'Java Application' launch configuration to use to create a runnable JAR.

**Launch configuration:**

Client - bench

**Export destination:**

C:\Users\SLMT\Desktop\benchmark\client.jar

**Library handling:**

- ☒ Extract required libraries into generated JAR
- ☐ Package required libraries into generated JAR
- ☐ Copy required libraries into a sub-folder next to the generated JAR

☐ Save as ANT script

**ANT script location:** D:\Courses\nthu-intro-dbms-2020\db20-assignment-4

Buttons: < Back, Next >, Finish, Cancel



# Setup Working Directory

- The next step is to setup you working directory.
- Contents

## – Server

- server.jar
- Properties
- Scripts



benchmark > server				
名稱	修改日期	類型	大小	
logging.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	3 KB	
server.jar	2020/5/2 下午 03:49	Executable Jar File	1,932 KB	
vanillabench.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	4 KB	
vanilladb.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	7 KB	

## – Client

- client.jar
- Properties
- Scripts



benchmark > client				
名稱	修改日期	類型	大小	
client.jar	2020/5/2 下午 03:49	Executable Jar File	1,932 KB	
logging.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	3 KB	
vanillabench.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	4 KB	
vanilladb.properties	2020/5/2 下午 03:45	PROPERTIES 檔案	7 KB	

# Scripts

- Now we are going to write scripts for running client and servers
- Scripts
  - Server
    - server.sh
    - copy-db.sh/reset-db.sh
  - Client
    - client-load.sh
    - client-bench.sh

# Execution Scripts

- server.sh

```
java -Djava.util.logging.config.file=logging.properties -  
Dorg.vanilladb.bench.config.file=vanillabench.properties -  
Dorg.vanilladb.core.config.file=vanilladb.properties -jar server.jar [DB Name]
```

- client-load.sh

```
java -Djava.util.logging.config.file=logging.properties -  
Dorg.vanilladb.bench.config.file=vanillabench.properties -  
Dorg.vanilladb.core.config.file=vanilladb.properties -jar client.jar 1
```

- client-bench.sh

```
java -Djava.util.logging.config.file=logging.properties -  
Dorg.vanilladb.bench.config.file=vanillabench.properties -  
Dorg.vanilladb.core.config.file=vanilladb.properties -jar client.jar 2
```

# Backup Databases

- To ensure the consistency of experiments, we usually backup the database and reset it before each experiment.
- copy-db.sh

```
DB_DIR="[DB Path]"  
cp -r $DB_DIR $DB_DIR-backup
```

- reset-db.sh

```
DB_DIR="[DB Path]"  
rm -r $DB_DIR  
cp -r $DB_DIR-backup $DB_DIR
```

# The Workflow of Benchmarking (1/2)

## 1. Load DB

1. Setup properties
2. Run `server.sh`
3. Run `client-load.sh`
4. Wait for loading
5. Shut down the server (by stopping the script)
6. Run `copy-db.sh`

# The Workflow of Benchmarking (2/2)

## 2. Benchmark

1. Setup properties
2. Run `reset-db.sh`
3. Run `server.sh`
4. Run `client-bench.sh`
5. Wait for benchmarking
6. Shut down the server (by stopping the script)