

ARIES

DataLab

Introduction to Database Systems

2022 Spring

What we expected

memory pages

page 20



Tx1 commit

Tx2 abort

page 20



Flush



- In page 20 :
 - Tx1 is a winner tx
 - Tx2 is a loser tx

However...

- Steal
 - Due to buffer management, dirty pages may be flushed to disk before txs commit
 - The changes made by **loser** txs must UNDO
- No Force
 - Due to performance reason, dirty pages won't get flushed immediately after txs commit
 - The changes made by **winner** txs must REDO

Logs in ARIES

Physical Log Record

- Record format :
 - Set Value Record
<Op Code, txNum, fileName, blockNum, offset, sqlType, **oldVal**, **newVal** >
 - Index Page Insert/Delete Record :
<Op Code, txNum, fileName, blockNum, insertSlot, insertKey, insertRidBlkNum, insertRidId>
- REDO :
 - Apply newVal to the page
- UNDO :
 - Apply oldVal to the page
 - Append its ***Compensation Log Record***

Compensation Log Record

- A **CLR** describes the actions taken to undo the actions of a previous update record.
- CLRs are added to log like any other record.
- Only need to **Redo** CLRs.
- It has all the fields of an update log record plus the **undoNext** pointer (the next-to-be-undone LSN).

Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

Redo



Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

Redo



Undo

[4]<SetValClr , 1, Page 20 , 3 , 2 > // Append Undo [3]

Redo log

Crash Again !

Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr , 1, Page 20 , 3 , 2 >

Crash Again !

Redo



Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4] <SetValClr , 1, Page 20 , 3 , 2 >

Crash Again !

[5] <SetValClr , 1, Page 20 , 2 , 3 > // Append Undo { Undo [3]
Redo log } Redo log

Crash Again !

Redo



Undo

Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr , 1, Page 20 , 3 , 2 >

Crash Again !

[5]<SetValClr , 1, Page 20 , 2 , 3 >

Crash Again !

Redo



Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr , 1, Page 20 , 3 , 2 >

Crash Again !

[5]<SetValClr , 1, Page 20 , 2 , 3 >

Crash Again !

[6]<SetValClr , 1, Page 20 , 3 , 2 >

// Append Undo { Undo { Undo [3] Redo log } Redo log} Redo log

...

Redo



Undo

Why CLR is Redo Only ?

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4] <SetValClr , 1, Page 20 , 3 , 2 >


Crash Again !

Redo



Undo

Why CLR is Redo Only ?



How do we know where Undo should start?

Why CLR Needs UndoNext?

- UndoNext helps skip logs which have been Undone by Redo

[0] <Start 1>

[1] <SetVal , 1 , Page 20 , 0 , 1>

[2] <SetVal , 1 , Page 20 , 1 , 2>

[3] <SetVal , 1 , Page 20 , 2 , 3>

Crash Here !

[4]<SetValClr 1, Page 20 , 3 , 2 , [3] > // Append Undo [3] Redo log

Crash Again !

[5]<SetValClr 1, Page 20 , 2 , 1 , [2] > // Append Undo [2] Redo log

[6]<SetValClr 1, Page 20 , 1 , 0 , [1] >

Logical Log Record

- Record format :
 - Logical - Start Record
<OP Code, txNum>
 - Record File Insert/Delete End Record :
<Op Code, txNum, fileName, blockNum, slotId , logicalStartLSN>
 - Index Insert/Delete End Record :
<Op Code, txNum, tblName, fldName, searchKey, recordBlockNum, recordSlotId , logicalStartLSN>
- REDO :
 - Do nothing
- UNDO :
 - Undo **completed** logical log **logically**
 - Undo **partial** logical log **physically**
 - Append **Logical Abort** log record

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [1] >

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [1] >

Logical
operations

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page Insert , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [1] >

Physical
operations

Logical
operations

Crash Here !

Rollback a completed Logical Log Record

[0] <Start 1>

[1] <LogicalStart, 1 >

[2] <Index Page **Insert** , 1 , ... >

[3] <SetVal , 1 , Page 2 , 1 , 2>

[4] <SetVal , 1 , Page 20 , 2 , 3>

[5] <Record File Insert End , 1, ... , [1] >

Crash Here !

[6] <Start 2 >

[7] <LogicalStart, 2 >

[8] <Index Page **Delete** , 2 , ... >

[9] <SetVal , 2 , Page 2 , 2 , 1>

[10] <SetVal , 2 , Page 20 , 3 , 2>

[11] <Record File Delete End , 2, ... , [7]>

[12] <Logical Abort 1 , [1]>


Physical
operations

Logical
operations

Recovery Phases of ARIES


Recovery Phases of ARIES

- Analysis Phase
 - Find the earliest possible start point of dirty page
 - Find loser txs




Recovery Phases of ARIES

- Analysis Phase
 - Find the earliest possibly start point of dirty page
 - Find loser txs
- REDO Phase
 - Repeat history (*both* winner and loser changes)
 - Recovery exact page status when the failure occurred

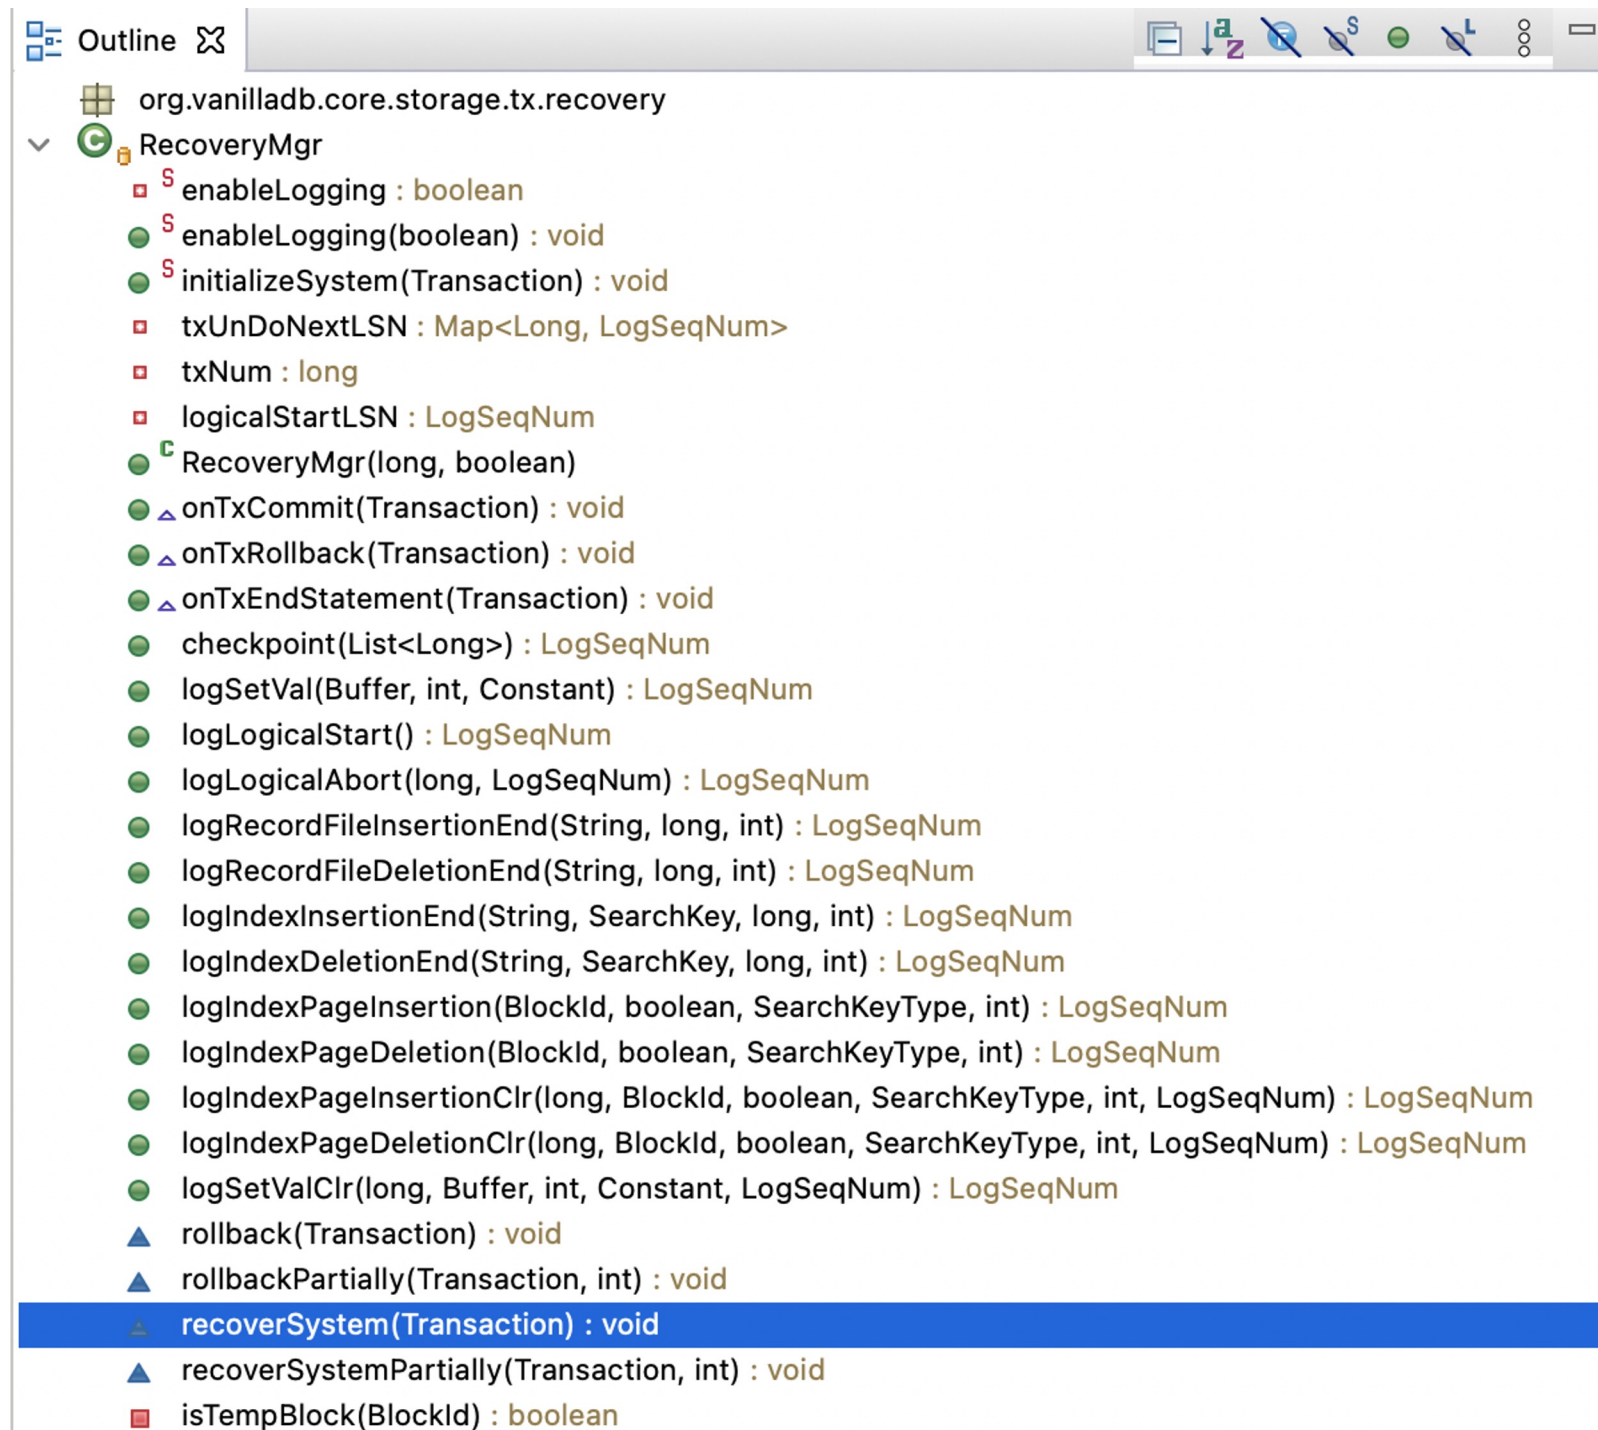


Recovery Phases of ARIES

- Analysis Phase
 - Find the earliest possibly start point of dirty page
 - Find loser txs
- REDO Phase
 - Repeat history (*both* winner and loser changes)
 - Recovery exact page status when the failure occurred
- UNDO Phase
 - Rollback *loser* txs changes




Recovery in VanillaDB



The screenshot shows an IDE window with the 'Outline' tab selected. The package path is 'org.vanilladb.core.storage.tx.recovery'. The 'RecoveryMgr' class is expanded, showing its methods and fields. The methods are listed with their return types and parameters. The 'recoverSystem(Transaction) : void' method is highlighted in blue.

- org.vanilladb.core.storage.tx.recovery
- RecoveryMgr
 - enableLogging : boolean
 - enableLogging(boolean) : void
 - initializeSystem(Transaction) : void
 - txUndoNextLSN : Map<Long, LogSeqNum>
 - txNum : long
 - logicalStartLSN : LogSeqNum
 - RecoveryMgr(long, boolean)
 - onTxCommit(Transaction) : void
 - onTxRollback(Transaction) : void
 - onTxEndStatement(Transaction) : void
 - checkpoint(List<Long>) : LogSeqNum
 - logSetVal(Buffer, int, Constant) : LogSeqNum
 - logLogicalStart() : LogSeqNum
 - logLogicalAbort(long, LogSeqNum) : LogSeqNum
 - logRecordFileInsertionEnd(String, long, int) : LogSeqNum
 - logRecordFileDeletionEnd(String, long, int) : LogSeqNum
 - logIndexInsertionEnd(String, SearchKey, long, int) : LogSeqNum
 - logIndexDeletionEnd(String, SearchKey, long, int) : LogSeqNum
 - logIndexPageInsertion(BlockId, boolean, SearchKeyType, int) : LogSeqNum
 - logIndexPageDeletion(BlockId, boolean, SearchKeyType, int) : LogSeqNum
 - logIndexPageInsertionClr(long, BlockId, boolean, SearchKeyType, int, LogSeqNum) : LogSeqNum
 - logIndexPageDeletionClr(long, BlockId, boolean, SearchKeyType, int, LogSeqNum) : LogSeqNum
 - logSetValClr(long, Buffer, int, Constant, LogSeqNum) : LogSeqNum
 - rollback(Transaction) : void
 - rollbackPartially(Transaction, int) : void
 - recoverSystem(Transaction) : void**
 - recoverSystemPartially(Transaction, int) : void
 - isTempBlock(BlockId) : boolean

Analysis Phase




```
// analyze phase
while (iter.hasNext()) {
    LogRecord rec = iter.next();

    int op = rec.op();
    if (op == OP_CHECKPOINT) {
        // Since we flush all dirtyPage at checkpoint, therefore no need
        // to find the start record of active txNum
        txsOnCheckpointing = ((CheckpointRecord) rec).activeTxNums();
        for (long acTxn : txsOnCheckpointing) {
            // txNum give us info of possible unFinshedTxS,
            // Check if those weren't in finishedTxS, and add it to the
            // uncompletedTxS
            if (!finishedTxS.contains(acTxn))
                unCompletedTxS.add(acTxn);
        }
        // Start Redo From checkpoint
        break;
    }

    if (op == OP_COMMIT) {
        finishedTxS.add(rec.txNumber());
    } else if (op == OP_ROLLBACK) {
        finishedTxS.add(rec.txNumber());
    } else if (op == OP_START && !finishedTxS.contains(rec.txNumber())) {
        unCompletedTxS.add(rec.txNumber());
    }
}
```

Redo Phase



```
/*
 * redo phase: Repeating History
 */
while (iter.hasPrevious()) {
    LogRecord rec = iter.previous();
    rec.redo(tx);
}
```

Undo Phase

```
/*
 * undo phase: undo all actions performed by the active txs during last
 * crash
 */
while (iter.hasNext()) {
    LogRecord rec = iter.next();

    int op = rec.op();
    if (!unCompletedTxs.contains(rec.txNumber()) || op == OP_COMMIT || op == OP_ROLLBACK)
        continue;
    /*
     * Use UndoNextLSN to skip unnecessary physical record which have
     * been redo its undo by CLR or records have been rolled back
     */

    if (txUndoNextLSN.containsKey(rec.txNumber())) {
        if (txUndoNextLSN.get(rec.txNumber()).compareTo(rec.getLSN()) != 1)
            continue;
    }
    if (op == OP_START)
        unCompletedTxs.remove(rec.txNumber());
    else if (rec instanceof LogicalEndRecord) {

        // Undo this Logical operation;
        rec.undo(tx);

        LogSeqNum logicalStartLSN = ((LogicalEndRecord) rec).getlogicalStartLSN();
        /*
         * Save the Logical Start LSN to skip the log records between
         * the end record and the start record
         */
        txUndoNextLSN.put(rec.txNumber(), logicalStartLSN);
    } else if (rec instanceof CompesationLogRecord) {

        LogSeqNum undoNextLSN = ((CompesationLogRecord) rec).getUndoNextLSN();
        /*
         * Save the UndoNext LSN to skip the records have been rolled
         * back
         */
        txUndoNextLSN.put(rec.txNumber(), undoNextLSN);
    } else
        rec.undo(tx);

    if (unCompletedTxs.size() == 0)
        break;
}
```

Reference

- ARIES: a transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging