

NewSQL and Determinism

Shan-Hung Wu and DataLab
CS, NTHU

Outline

- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

Outline

- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

The evolving database landscape

451 Research

Relational

Analytic

Hadoop Piccolo Teradata Aster Netezza ParAccel SAP Sybase IQ
 Hadapt Infobright EMC Greenplum IBM InfoSphere
 HPCC RainStor Teradata Calpont Action VectorWise HP Vertica

Non-relational

NoSQL

MarkLogic

DataStax Enterprise
 Castle Acunu

Versant

Citrusleaf
 BerkeleyDB Cassandra HBase

McObject

Oracle NoSQL
 RethinkDB
 HandlerSocket*

Progress

LevelDB
 Redis
 Membrain
 Voldemort

Objectivity

Redis-to-go
 SimpleDB
 DynamoDB
 Iris Mongo Mongo
 Couch Lab HQ
 Couchbase RavenDB
 MongoDB CouchDB

Graph

Neo4J
 InfiniteGraph
 OrientDB
 DEX
 NuvolaBase

-as-a-Service

Big tables

App Engine

Datastore

Key value

Document

Lotus Notes

Operational

Starcounter

InterSystems

SAP HANA

Oracle

Percona

IBM DB2

MariaDB

SkySQL

MySQL

PostgreSQL

SQL Server

-as-a-Service

FathomDB
 Amazon RDS
 Database.com
 Postgres Plus Cloud
 ClearDB
 Rackspace MySQL Cloud
 Google Cloud SQL
 SQL Azure

Action Ingres

EnterpriseDB

SAP Sybase ASE

NewSQL

-as-a-Service

StormDB

Xeround

Storage engines

Tokutek

ScaleDB

Zimory

Scale

Continuent

MySQL Cluster

Galera

CodeFutures

ScaleBase

Clustering/sharding

New databases

NuoDB VoltDB

MemSQL JustOneDB SQLFire

Drizzle Akiban Translattice

SchoonerSQL Clustrix

GenieDB ScaleArc ParElastic

ScaleDB Zimory Scale Continuent

MySQL Cluster Galera CodeFutures

ScaleBase Clustering/sharding

Shall we give up the relational model?

What To Lose?

- Expressive model
 - Normal forms to avoid insert/write/delete anomalies
- Flexible queries
 - Join, group-by, etc.
- Txns and ACID across tables

Workarounds

- No expressive model
 - Your application/data must fit
- No flexible queries
 - Issue multiple queries and post-process them
- No tx and ACID across key-value pairs/documents/entity groups
 - Code around to access data “smartly” to simulate cross-pair txs

“We believe it is better to have application programmers deal with performance problems due to overuse of transactions as bottlenecks arise, rather than always coding around the lack of transactions”

- an ironic to NoSQL when Google releases Spanner

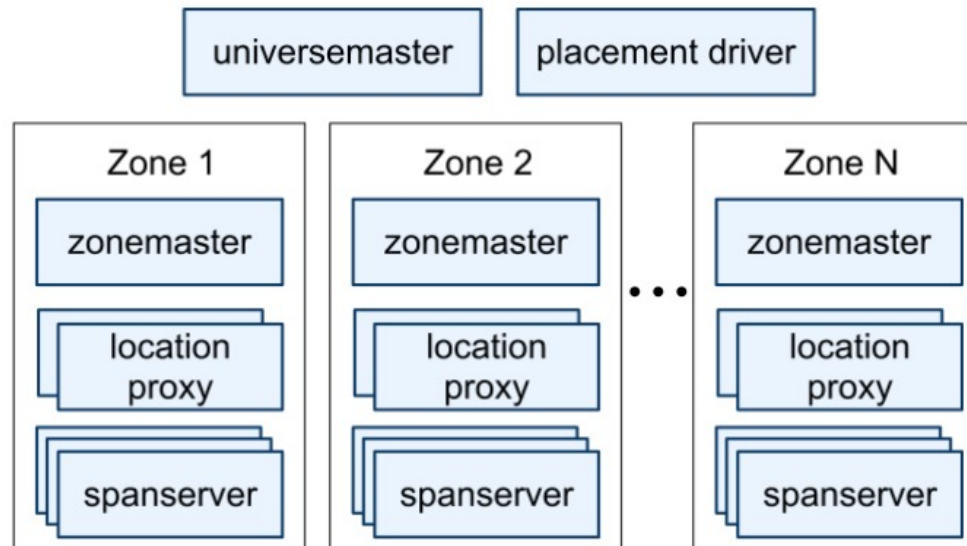
Is full DB functions + SAE possible?

Google Spanner

- Descendant of Bigtable, successor to Megastore
- Properties
 - Globally distributed
 - Synchronous cross-datacenter replication (with Paxos)
 - Transparent sharding, data movement
 - General transactions
 - Multiple reads followed by a single atomic write
 - Local or cross-machine (using 2PC)
 - Snapshot reads

Spanner Server Organization

- Zones are the unit of physical isolation
- Zonemaster assigns data to spanservers
- Spanserver serves data (100~1000 tablets)




Data Model

- Spanner is semi-relational model
 - Every table has primary key (unique name of row)
- The *hierarchy* abstraction (assigned by client)
 - INTERLEAVE IN allows clients to describe the locality relationships that exist between multiple tables

```
CREATE TABLE Users {  
  uid INT64 NOT NULL, email STRING  
} PRIMARY KEY (uid), DIRECTORY;
```

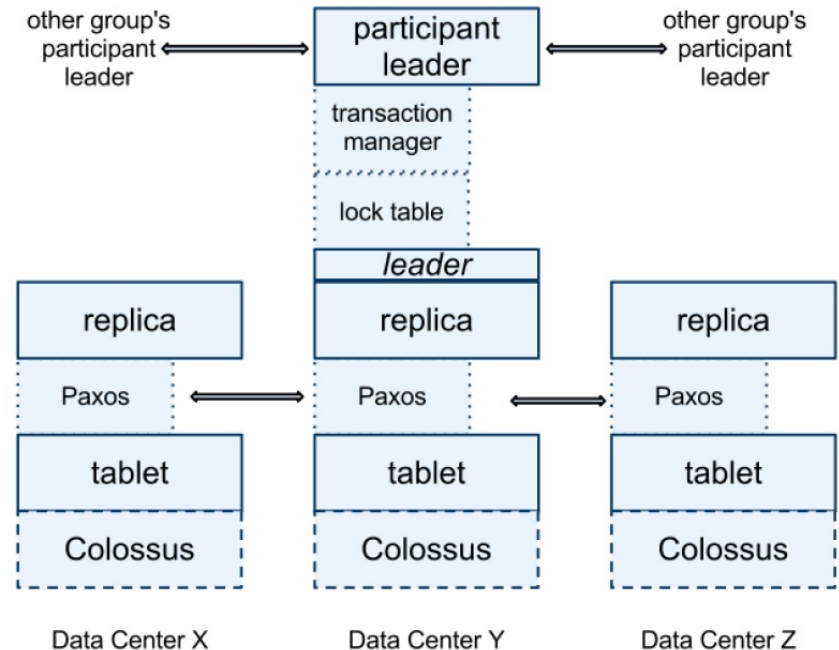
```
CREATE TABLE Albums {  
  uid INT64 NOT NULL, aid INT64 NOT NULL,  
  name STRING  
} PRIMARY KEY (uid, aid),  
  INTERLEAVE IN PARENT Users ON DELETE CASCADE;
```



- Same hierarchy → same partition


Transaction Execution

- Writes are buffered at client until commit
- Commit: use Paxos to determine global serial order
- Multiple Paxos groups
 - Same hierarchy → same group
- Single Paxos group txn
 - bypass txn mgr
- Multi Paxos group txn
 - 2PC done by txn mgr



Directories

- Directories are the unit of data movement between Paxos groups
 - shed load
 - frequently accessed together
- Row's key in Spanner tablet may not be contiguous (unlike Bigtable)

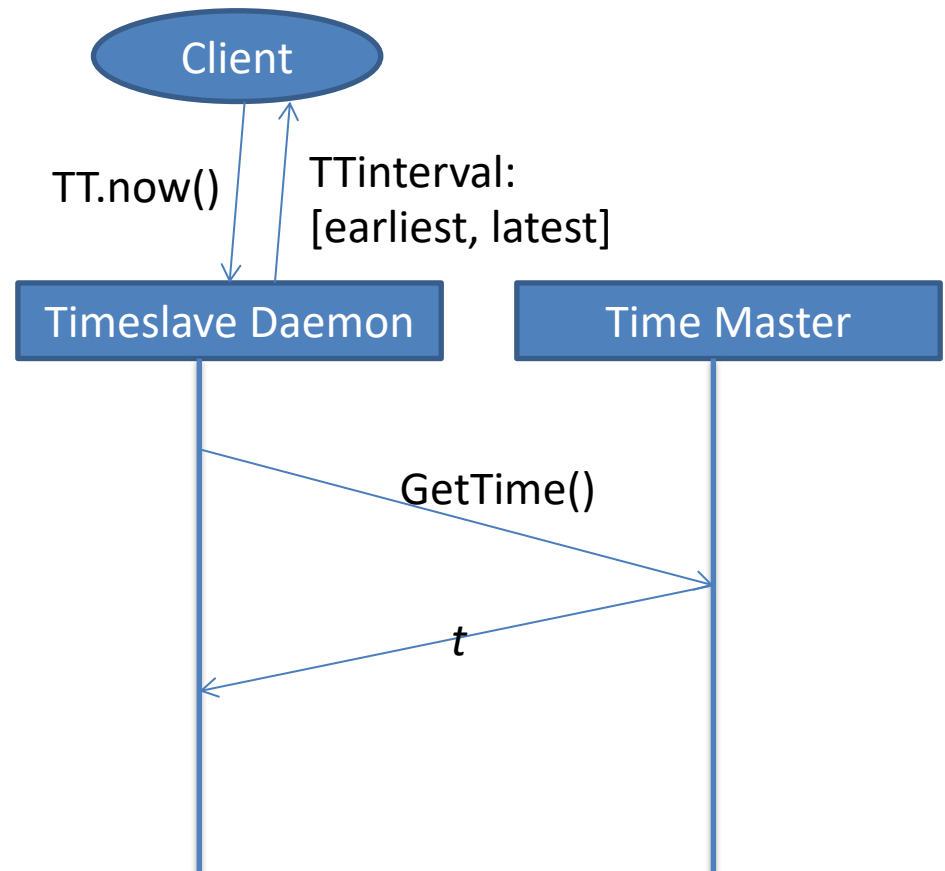


Isolation & Consistency

- Supported txns
 - Read-Write txn (need locks)
 - Read-Only txn (up-to-date timestamp)
 - Snapshot Read
- MVCC-based
- Consistent read/write at global scale
- Every r/w-txn has commit timestamp
 - Reflects the serialization order
- Snapshot reads are not blocked by writes (at the cost of snapshot isolation)

True Time

- GPS + atomic clock
- $t_{abs}(e)$: absolute time of an event e
- TrueTime guarantees that for an invocation $tt = TT.now()$, $tt.earliest \leq t_{abs}(e_{now}) \leq tt.latest$, *where e_{now} is the invocation event*
- *Error bound ϵ*
 - half of the interval's width
 - $1 \sim 7\ ms$

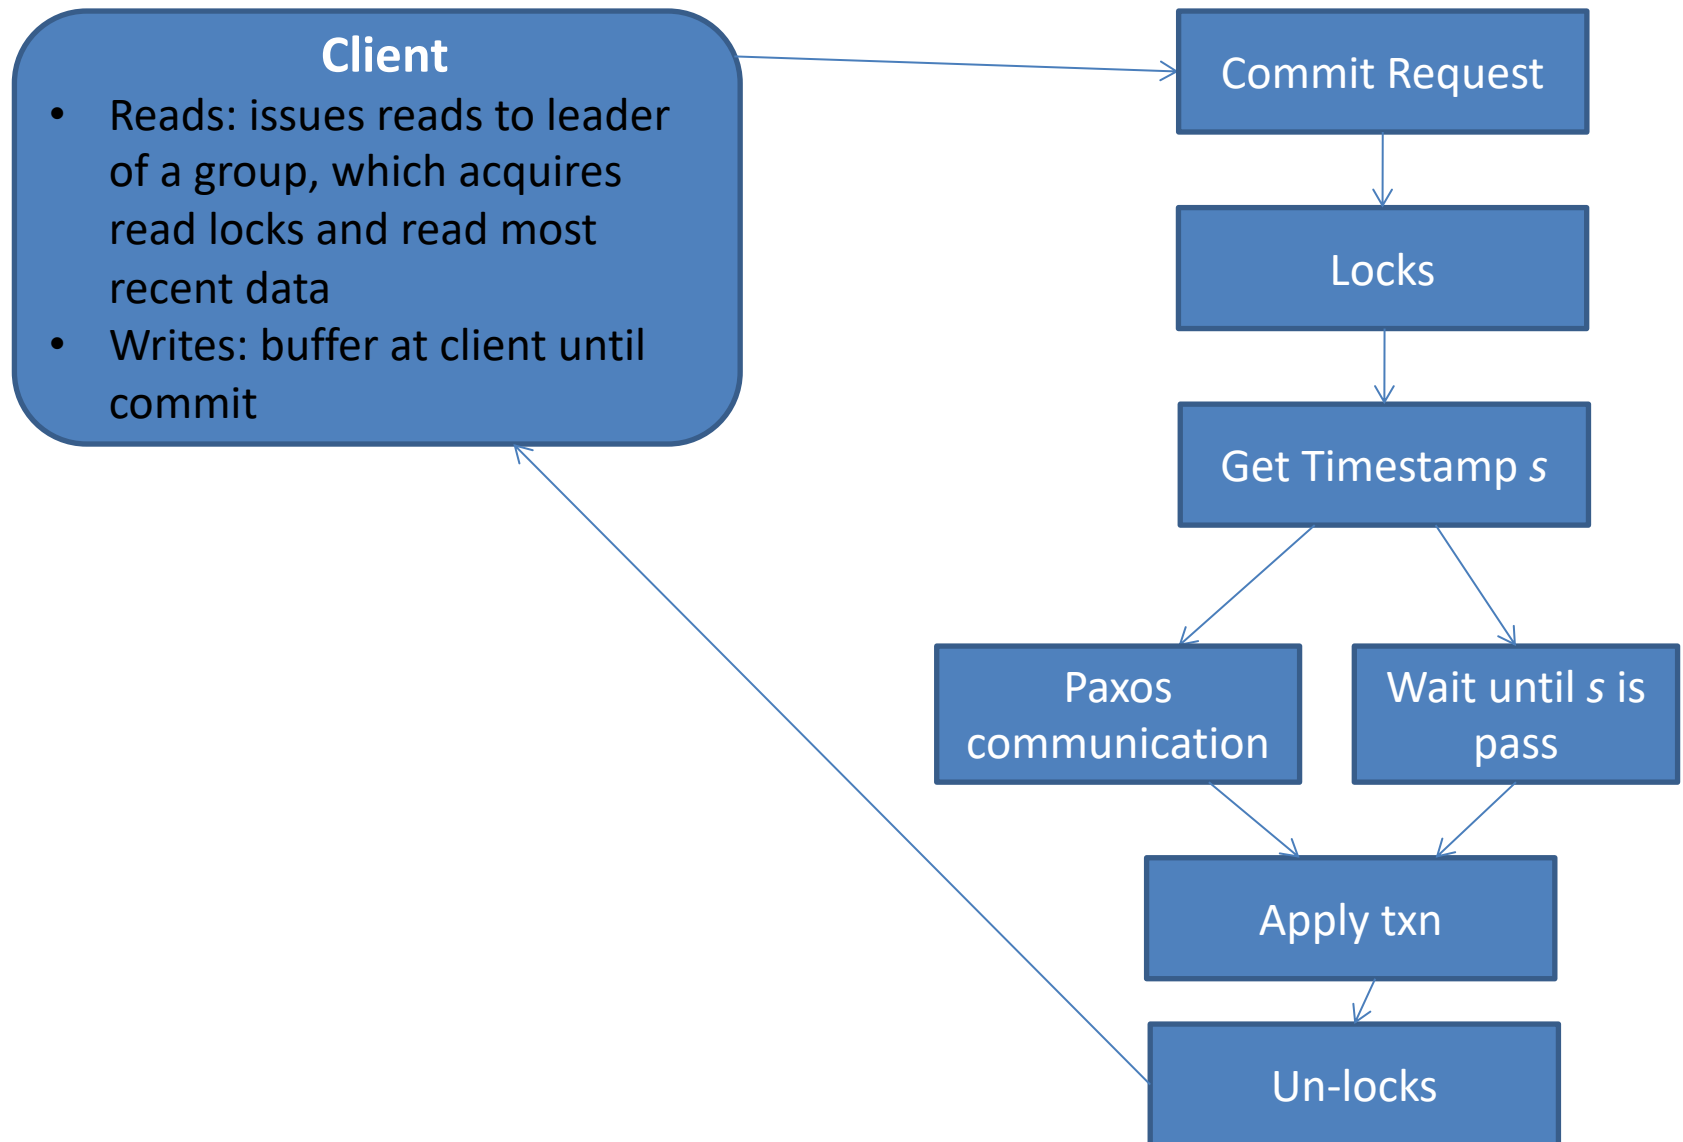


Timestamp Management

- Spanner enforces the following external consistency invariant
 - If the start of a transaction T2 occurs after the commit of a transaction T1 , then the commit timestamp of T2 must be greater than the commit timestamp of T1
 - Commit timestamp of a transaction Ti by s_i

$$t_{abs}(e_1^{commit}) < t_{abs}(e_2^{start}) \Rightarrow s_1 < s_2$$

Read-Write Txn



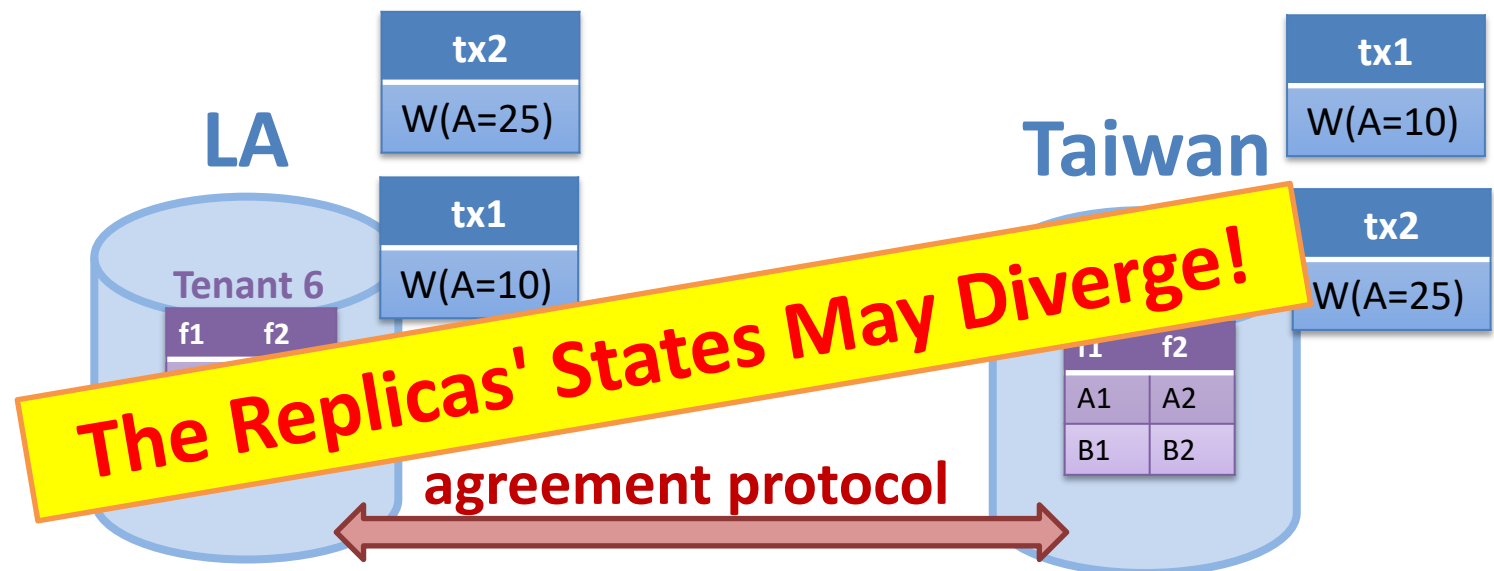
Outline

- The NewSQL landscape
- **Determinism for SAE**
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

What is a deterministic DDBMS?


What Is A Non-Deterministic DDBMS?

- Given Tx_1 and Tx_2 , traditional DBMS guarantees **some** serial execution: $Tx_1 \rightarrow Tx_2$, or $Tx_2 \rightarrow Tx_1$
- Given a collection of requests/txs, DDBMS leaves the data (outcome) non-deterministically due to
 - Delayed requests, CC (lock granting), deadlock handling, buffer pinning, threading, etc.



Deterministic DDBMS

- Given a collection of **ordered** requests/txs, if we can ensure that
 - conflicting txs are executed on each replica following that order
 - each tx is always run to completion (commit or abort by tx logic), even failure happens
- Then data in all replicas will be in same state, i.e., **determinism**



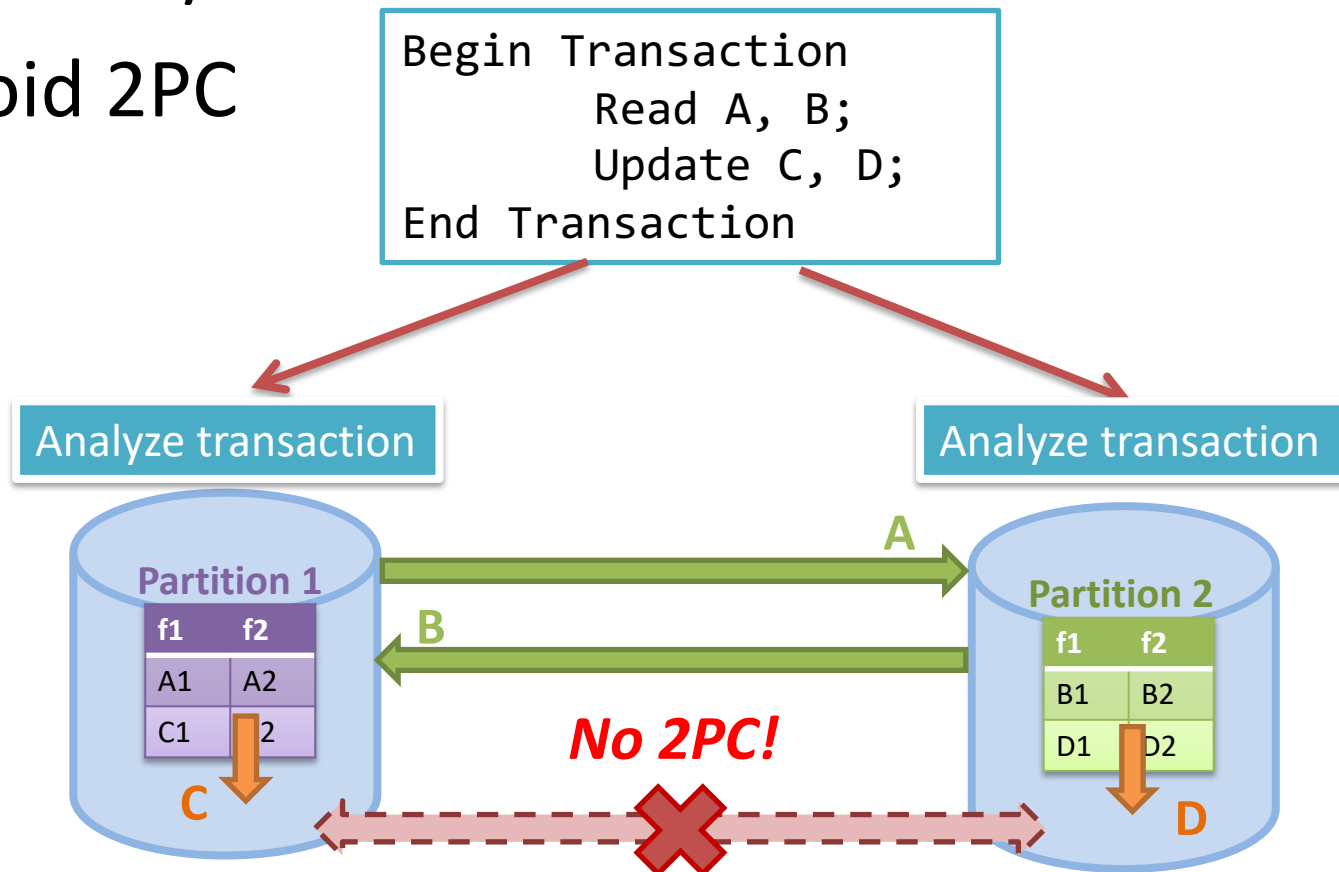
Wait, how about the delay in total-ordering?

Total-Ordering Cost

- Long tx latency
 - Hundreds of milliseconds
- However, the delay of total-ordering does *not* count into contention footprint
- No hurt on throughput!


Scalability via Push-based Dist. Txns.

- Push readings proactively
- Parallel I/Os
- Avoid 2PC



T-Part

- Better tx dispatching → forward pushes




Elasticity

- Based on currently workload, we need to decide:
 1. How to re-partition data
 2. How to migrate data?

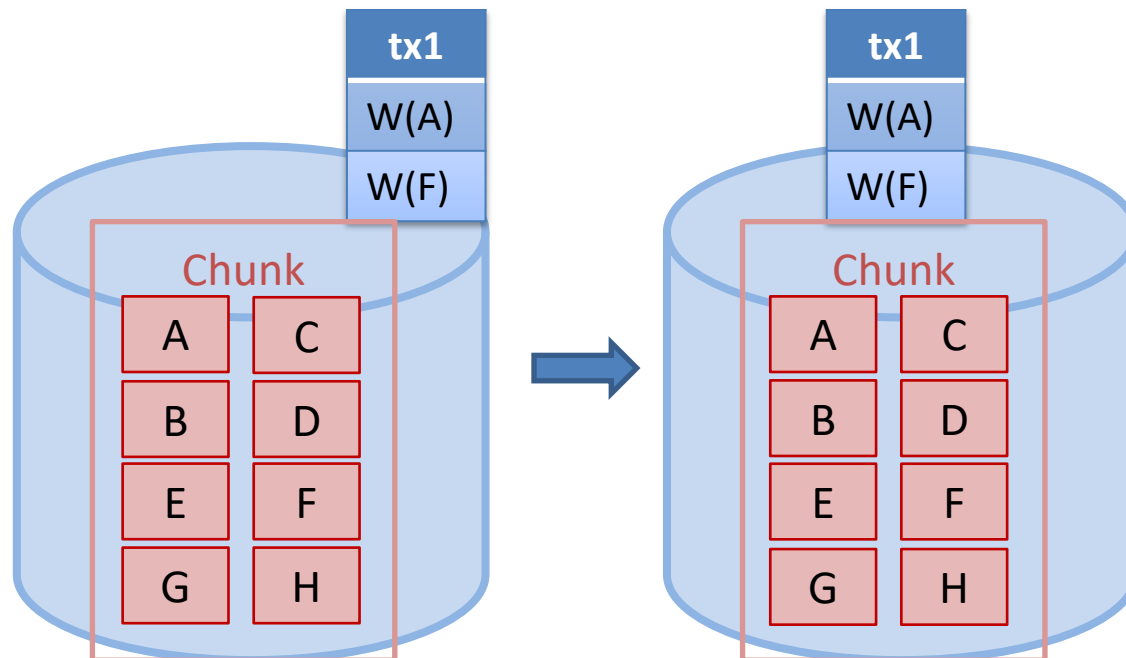
Re-Partitioning

- Data chunking
 - User-specified (e.g., Spanner hierarchy), or
 - System generated (e.g., [consistent hashing](#))
- Master server for load monitoring
 - Practical for load balancing in a single datacenter




Migration *without* Determinism

- Txs that access data being migrated are blocked
- User-aware down time due to slow migration process




Migration *with* Determinism

- Replication is *cheap* in deterministic DBMSs
- To migrate data from a source node to a new node, the source node *replicates* data by executing txs in *both* nodes




Foreground Pushing

- The source node *pushes* the data that txs need to access on the new node
 - If some of fetched records also need to be migrated, the new node will store those records




Background Pushing

- The rest of records which are not pushed by txs can be pushed by a background *async-pushing* thread
- Then, the source node deletes the migrated records



Migration vs. Crabbing

- Client served by **any** node running faster
 - Source node in the beginning; dest. node at last
- Migration delay imperceptible



How to implement a deterministic
DDBMS?

Outline

- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

Outline

- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

Implementing Deterministic DDBMS

- Requests need to be ordered
- A deterministic DBMS engine for each replica

Ordering Requests

- A ***total ordering protocol*** needed
 - Replicas have consensus on the order of requests
- A topic of ***group communication***
- To be discussed later

The Sources of Non-Determinism

- Thread and process scheduling
- Concurrency control
- Buffer and cache management
- Hardware failures
- Variable network latency
- Deadlock resolution schemes

How to make a DBMS deterministic
(given the total-ordered requests)?

Outline

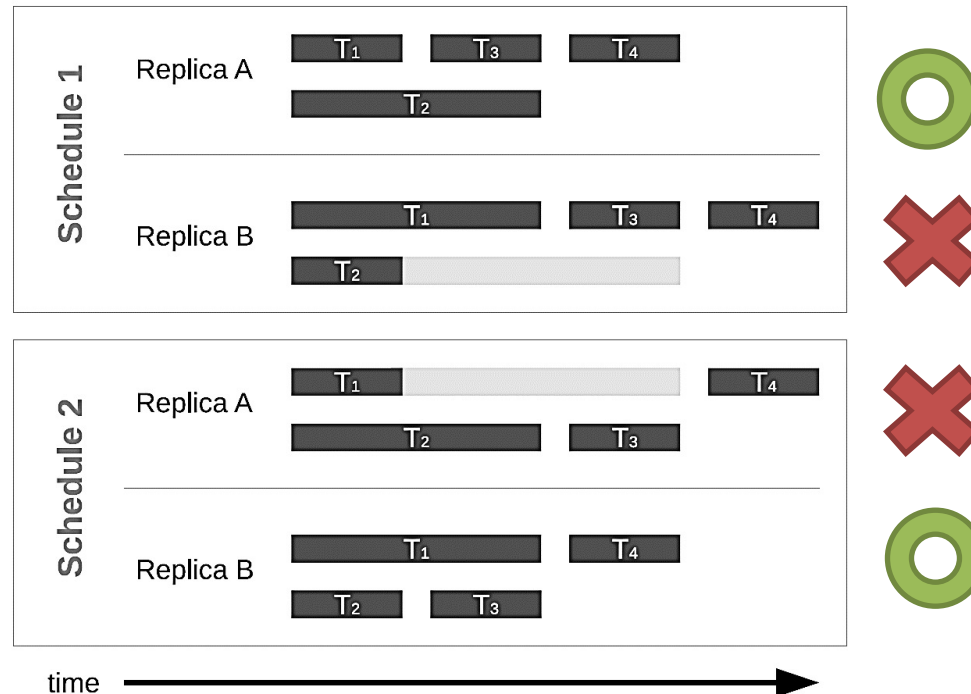
- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - **Ordered locking**
 - Dependent transactions
- Performance study

Single-Threaded Execution

- Execute the ordered requests/txs one-by one
- Good idea?
- Poor performance?
 - No pipelining between CPU and I/O or networking
 - Does not take the advantage of multi-core
 - The order of non-conflicting txs doesn't matter
- Walk-arounds: H-store


Fix-Ordered Multi-Thread Execution?

- T1, T2, and T3 do not conflict with each other
- T4 conflicts with all T1, T2, and T3



Dynamic Reordering is Better

- Schedule of conflicting txs are the same across all replicas
- Schedule of non-conflicting txs can be “*re-ordered*” automatically
 - Better pipelining to increase system utilization
 - Shorter latency



How? S2PL?

How to make S2PL deterministic?

A Case of Determinism

- Based on S2PL:
- ***Conservative***, ordered locking
- Each tx is always run to completion (commit or abort by tx logic), ***even failure happens***

Conservative, Ordered Locking

- Each tx has ID following the total order
- Requests *all* it locks *atomically* upon entering the system (conservative locking)
- For any pair of transactions T_i and T_j which both request locks on some record r , if $i < j$ then T_i must request its lock on r before T_j
- *No deadlock!*

Execute to Completion

- Each node logs “requests”
 - Operations are idempotent if a tx is executed multiple times at the same place of a schedule
- No system-decided rollbacks during recovery
 - Recovery manager rollbacks all operations of incomplete txs
 - But loads request logs and *replay* those incomplete txs following the total order

Wait, how to know which object to lock *before* executing a tx?

Outline

- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- Performance study

Dependent Transactions

- It isn't necessarily possible for every tx to know which object to lock before execution
- Example: *dependent txs*
 - SQL: UPDATE ... SET ... WHERE ...
 - x is a field; y are values of another field
 - Which to lock for y ?
 - Option1: the entire table where y belongs to

$U(x)$:
 $y \leftarrow read(x)$
 $write(y)$

Rewriting Dependent Txs

- Option 2: rewrite tx!

```
U(x) :  
  y ← read(x)  
  write(y)
```



Reconnaissance tx

```
U1(x) :  
  y ← read(x)  
  newtxnrequest(U2(x, y))
```

and

```
U2(x, y) :  
  y' ← read(x)  
  if (y' ≠ y)  
    newtxnrequest(U2(x, y'))  
    abort()  
  else  
    write(y)
```

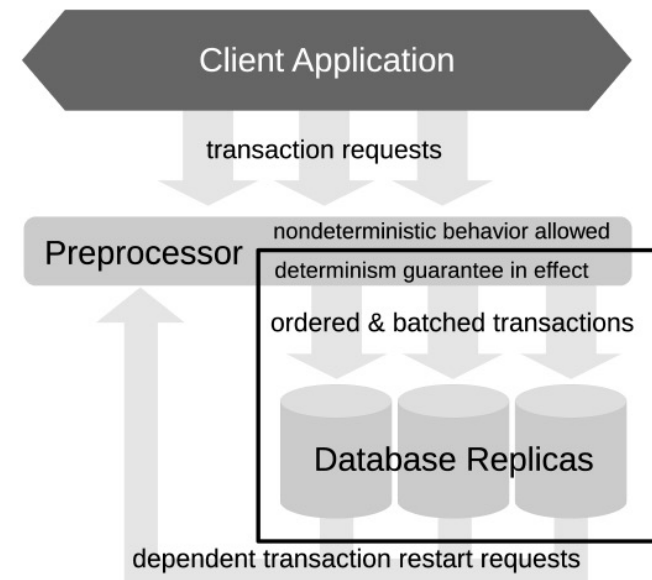


figure 2: Architecture of a deterministic DBMS.

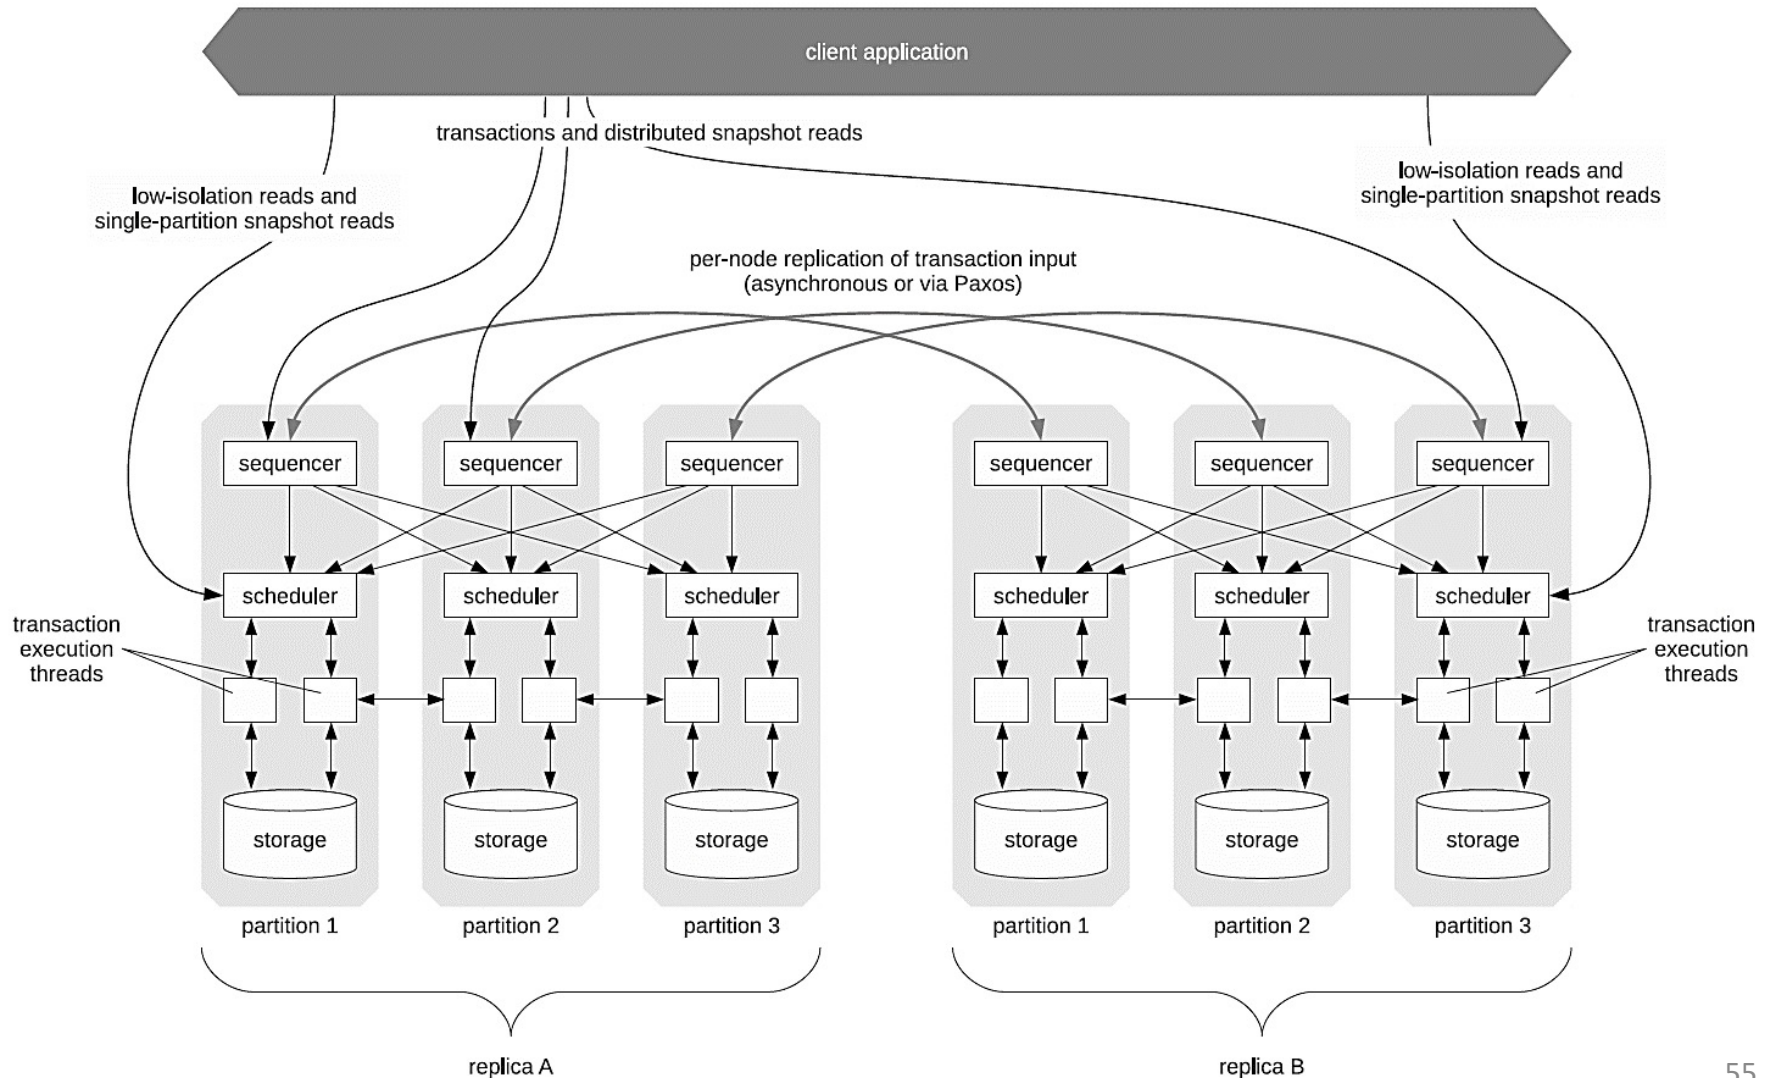
Rewriting Dependent Txs

- Pros?
 - Better concurrency for y (higher throughput)
- Overhead?
 - Increased latency
 - Higher abort rate if values in x are modified between $U1$ and $U2$
- Luckily, values in x may not be modified that often in practice
- Why?
 - Because otherwise you won't use x as a search field
 - x is usually (secondarily) indexed

Outline


- The NewSQL landscape
- Determinism for SAE
- Implementing a deterministic DDBMS
 - Ordering requests
 - Ordered locking
 - Dependent transactions
- **Performance study**

A Deterministic DBMS Prototype



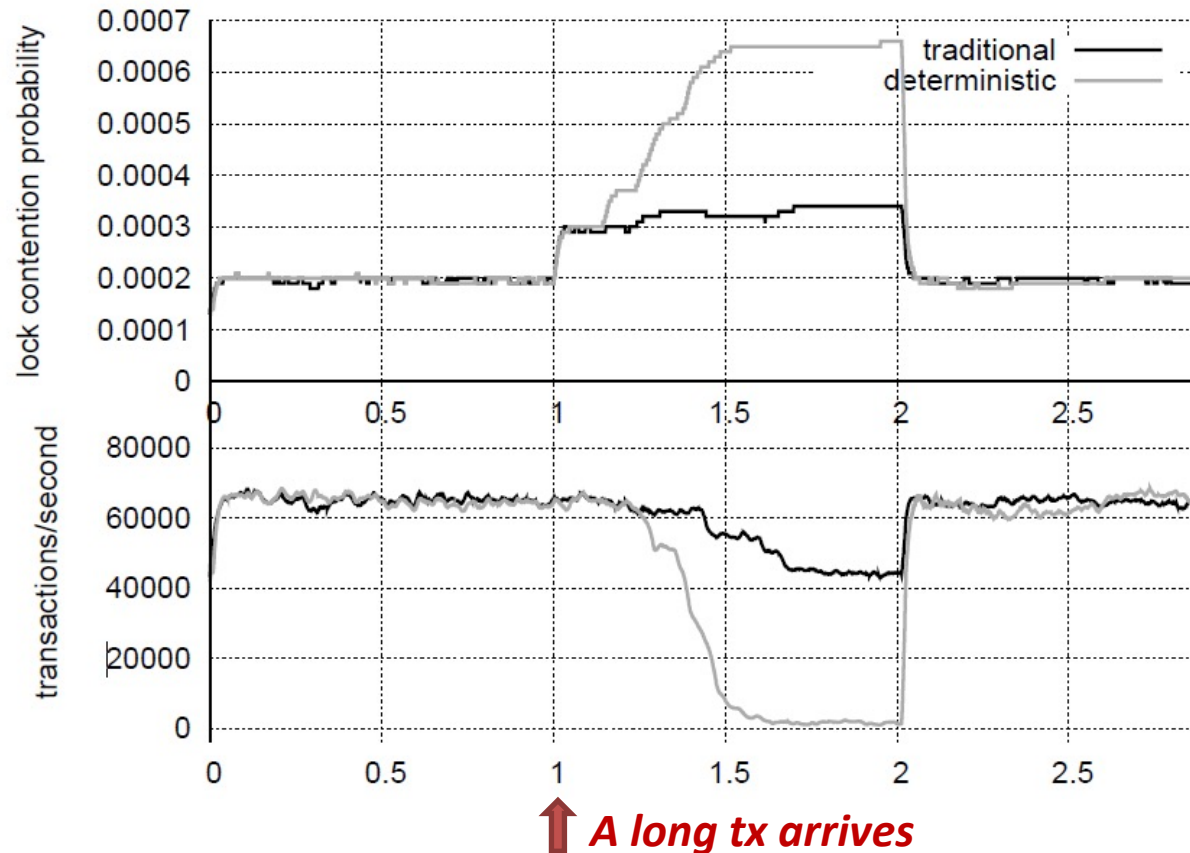
The Scalability of The Prototype

- Workload: TPC-C New Order Transactions



The Clogging Problem

- Workload: two txs conflict with probability 0.01%



Why?

Stalled Transaction and Clogging

- Clogging due to *stalled txs*
 - Following up txs conflicting with the long tx
- Why not an issue in traditional S2PL?
 - On-the-fly “re-ordering” between conflicting txs
 - E.g., with wound-wait deadlock prevention
- Conservative CC is suitable only for *workloads consisting of short txs*
 - Less-stalled txs
- In-memory DBMS (or alike) is a good fit
 - E.g., modern web deployments with Memcached between app server and DBMS

Review: Assumptions/Observations

- No ad-hoc txs
 - All txs are predefined (e.g., as stored procedures)
- No long txs
 - Minimize delay
 - Analytic jobs are moved to backend (e.g., GFS+MapReduce)
- Working set is small
- Locality

The Score Sheet

	DDBMS	Dynamo/Cassandra/Riak	Deterministic DDBMS
Data model	Relational	<i>Key-value</i>	Relational
Tx boundary	Unlimited	<i>Key-value</i>	Unlimited, <i>but not ad-hoc</i>
Consistency	W strong, <i>R eventual</i>	W strong, <i>R eventual</i>	WR strong
Latency	Low (local partition only)	Low	<i>High</i> (due to total ordering requests)
Throughput	High (<i>scale-up</i>)	High (scale-out)	High (scale-out, comm. delay doesn't impact CC)
Bottleneck/SPF	<i>Masters</i>	<i>Masters</i>	No
Consistency (F)	W strong, <i>R eventual</i>	<i>WR eventual</i>	WR strong
Availability (F)	<i>Read-only</i>	Read-write	Read-write
Data loss (F)	<i>Some</i>	Some	No
Reconciliation	Not needed	<i>Needed (F)</i>	Not needed
Elasticity	<i>No</i>	<i>Manual, blocking migration</i>	Yes, automatic

- High latency can be “hidden” in some applications
 - E.g., Gmail (with the aid of AJAX)

Final Remarks

- Cloud databases: a un-unified world
 - Common goals: **S**cale-out (not -up), high **A**vailability, and **E**lasticity
- It's not about how many DB products in the market you know
- It's about knowing the trade-offs
 - Better architecture design if you are in DB field
 - To choose the right product(s) as a DBMS user
- Full DB functions + SAE for is still an active research problem

References

- Matthew Aslett. What we talk about when we talk about NewSQL.
http://blogs.the451group.com/information_management/2011/04/06/what-we-talk-about-when-we-talk-about-newsql/
- Matthew Aslett. NoSQL, NewSQL and Beyond.
http://blogs.the451group.com/information_management/2011/04/15/nosql-newsql-and-beyond/
- A Thomson et al. The Case for Determinism in Database Systems. 2010.
- A Thomson et al. Calvin: Fast Distributed Transactions for Partitioned Database Systems. 2012.
- Towards an Elastic and Autonomic Multitenant Database, A. Elmore et al. NetDB 2011.
- Workload-aware database monitoring and consolidation, C. Curino et al. SIGMOD 2011.