

Vector Database Systems

Shan-Hung Wu and DataLab
CS, NTHU

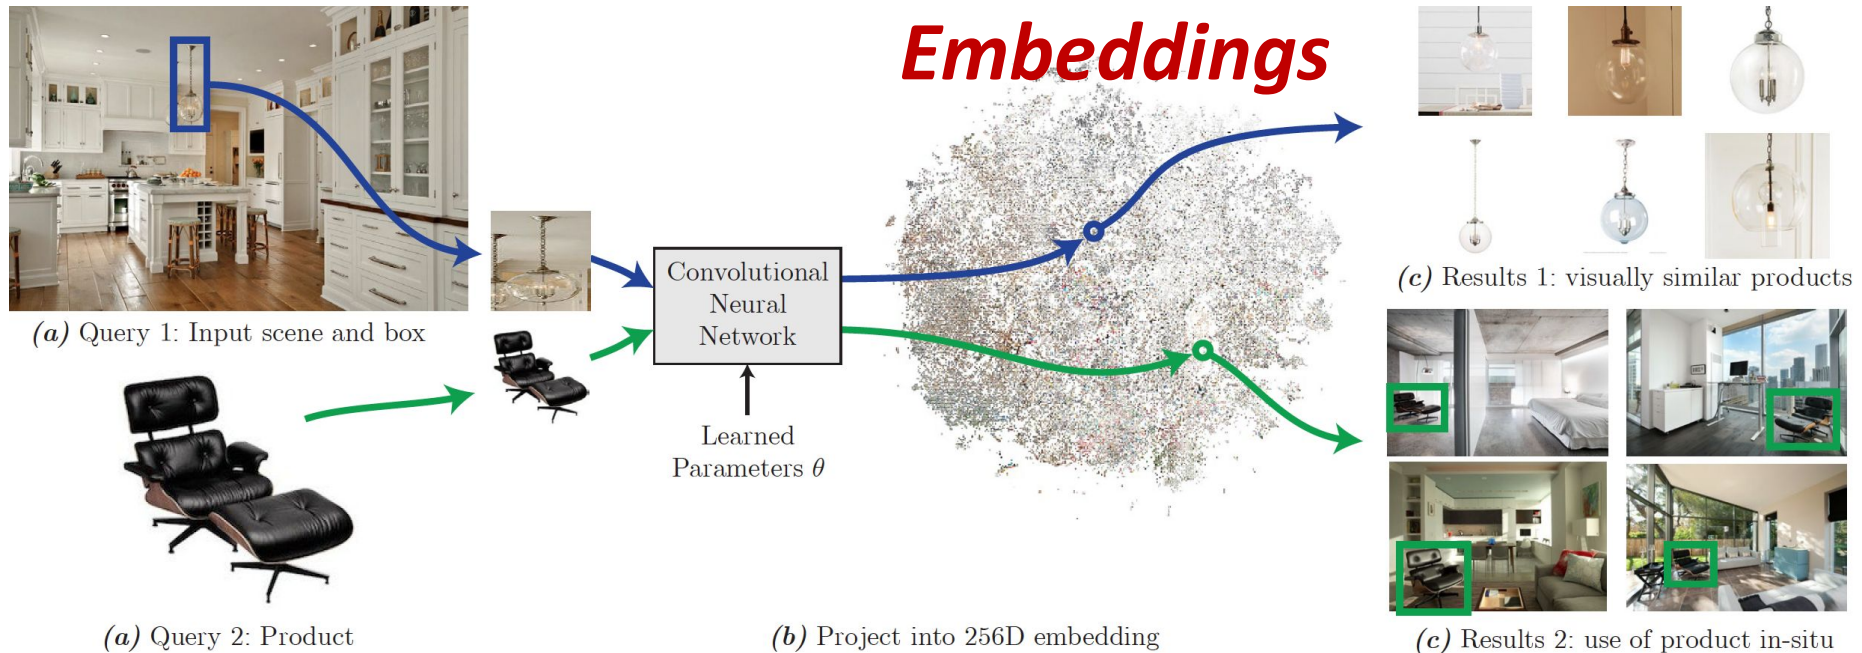
Outline

- Why Vector DBMS?
- AKNN Search Algorithms
- Challenges at System-level
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- Case study: Milvus (purpose-built)
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms


Outline

- **Why Vector DBMS?**
- AKNN Search Algorithms
- Challenges at System-level
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- Case study: Milvus (purpose-built)
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms

The Emerge of AI & Embeddings



- Used by search engines, recommender systems, personalized ads, etc.



How to store & search billions of embeddings?

Outline

- Why Vector DBMS?
- **AKNN Search Algorithms**
- Challenges at System-level
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- Case study: Milvus (purpose-built)
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms

Approximate K Nearest Neighbor (AKNN) Search

- Given a query vector q , find k vectors $V = \{v_1, v_2, \dots, v_k\}$ in storage that are approximately nearest to q
- Distance measure?
 - Euclidian distance, cosine similarity, etc.
- The higher *recall* the better
 - Let ground truth: V^*
 - Recall = $|V \cap V^*| / |V^*|$

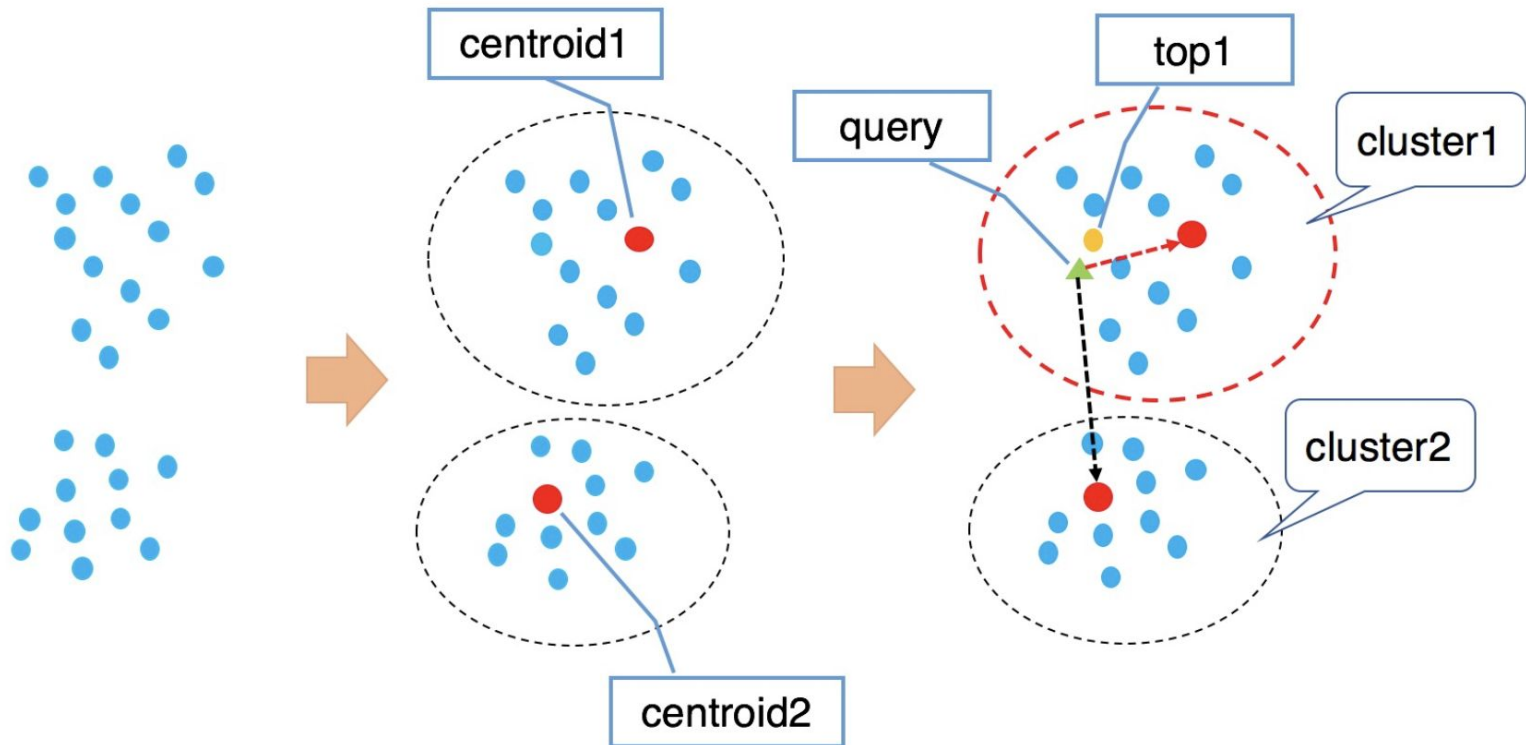
AKNN Algorithms

- Tree-based: KD-tree, R-tree
- Quantization-based: IVF_FLAT/SQ8/PQ
- Graph-based: HNSW, NSG, SSG
- Locality sensitive hashing (LSH)

AKNN Algorithms

- Tree-based: KD-tree, R-tree
 - Runs slowly on high-dimensional data
- ***Quantization-based***: IVF_FLAT/SQ8/PQ
 - ***High recall, codebooks are update-insensitive***
- Graph-based: HNSW, NSG, SSG
 - High recall, graph take time/space to maintain
- Locality sensitive hashing (LSH)
 - Low recall

IVF_FLAT/SQ8/PQ



- Search in each cluster: brut force (FLAT) vs. compressed (SQ8) vs. quantization of subvectors (PQ)

Outline

- Why Vector DBMS?
- AKNN Search Algorithms
- **Challenges at System-level**
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- Case study: Milvus (purpose-built)
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms

AKNN Libraries from AI Community

- Facebook **Faiss**, Microsoft **SPTAG**, Spotify **Annoy**, etc.
 - Implement various AKNN algorithms
- Pros: computation optimized
 - Support SIMD instructions (SSE, AVX, AVX2)
 - Faiss even supports GPU acceleration

AKNN Libraries from AI Community

- Facebook **Faiss**, Microsoft **SPTAG**, Spotify **Annoy**, etc.
 - Implement various AKNN algorithms
- **Cons:**
 - Assume memory storage only
 - No support for dynamic data (updates/deletes)
 - No attribute filtering (e.g., “ $100 < \text{price} < 200$ ”)

Outline

- Why Vector DBMS?
- AKNN Search Algorithms
- Challenges at System-level
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- Case study: Milvus (purpose-built)
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms

PHASE

- “PostgreSQL Ultra-High-Dimensional Approximate Nearest Neighbor Search Extension,” in SIGMOD’20
 - A PostgreSQL extension
 - Can be implemented in any System R-like DBMS
- Pros:
 - Supports disk storage
 - Supports dynamic data
 - Supports attribute filtering

Data Model

- Treats vectors as a *field* in a table
 - Type: `float_vector(d)`

- Index creation:

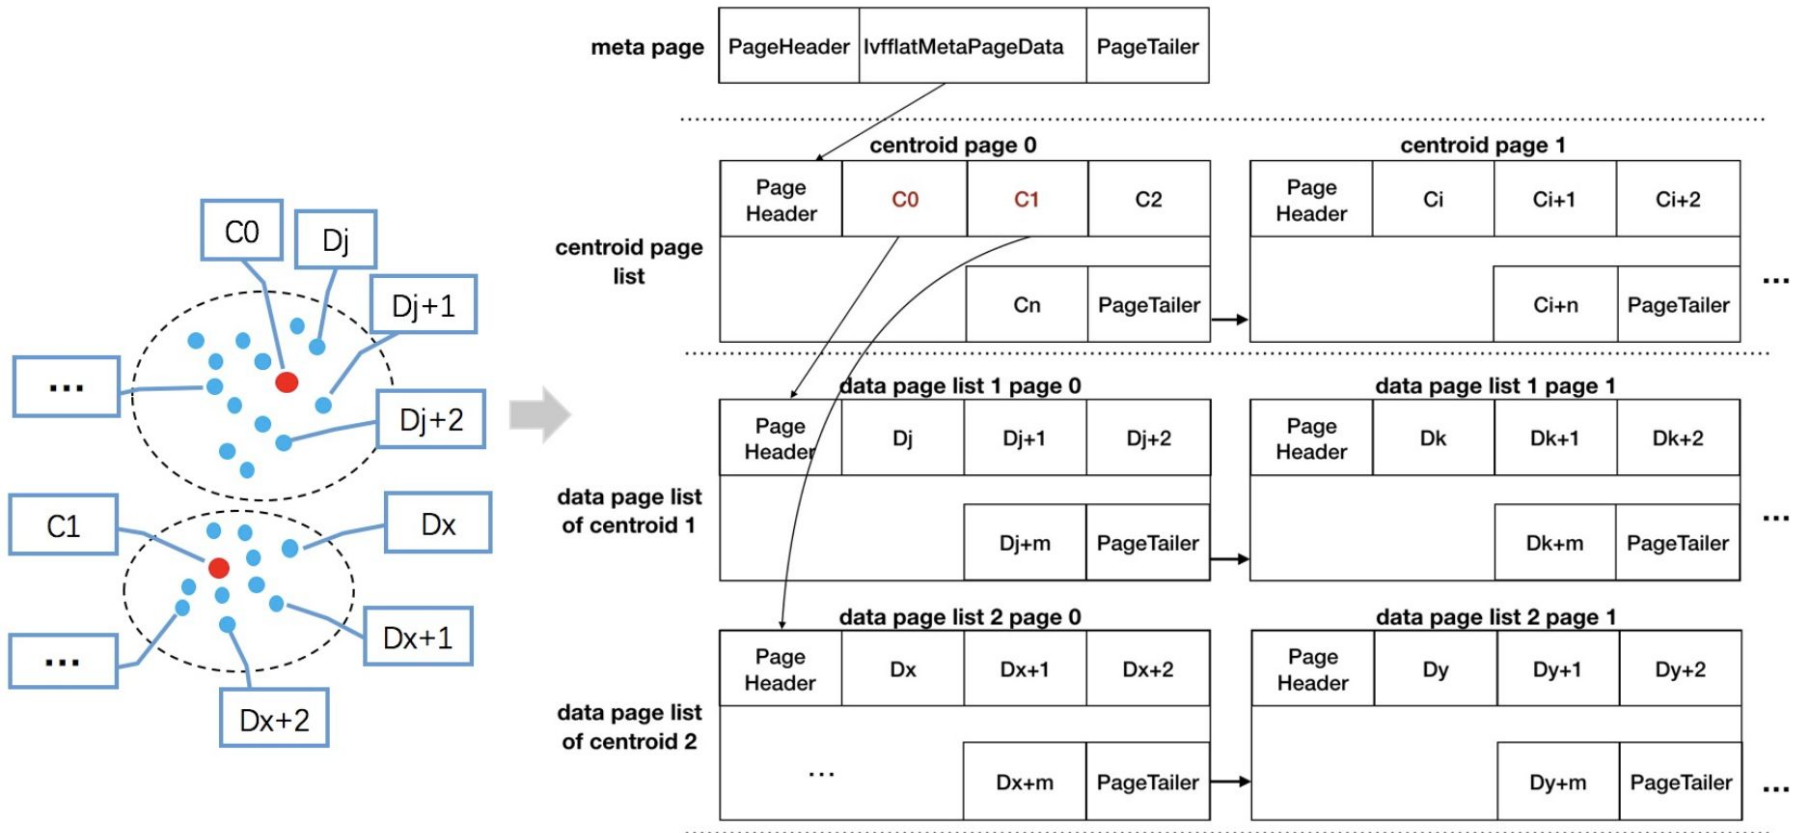
```
CREATE INDEX idx_text ON posts(text_vector)
USING ivf_flat;
```


Query Format

- AKNN query:

```
SELECT p.id,  
       p.text_vector <-> '...' AS dist  
FROM posts AS p  
ORDER BY dist ASC LIMIT 10;
```

Index Building (IVF_FLAT)



- Each page is the unit of buffering and searching

Index Update (IVF_FLAT)

- Do nothing if the data distribution does not change
- Otherwise, continue clustering for few iterations

Planning

```
SELECT p.id,  
       p.text_vector <-> '...' AS dist  
FROM posts AS p  
ORDER BY dist ASC LIMIT 10;
```

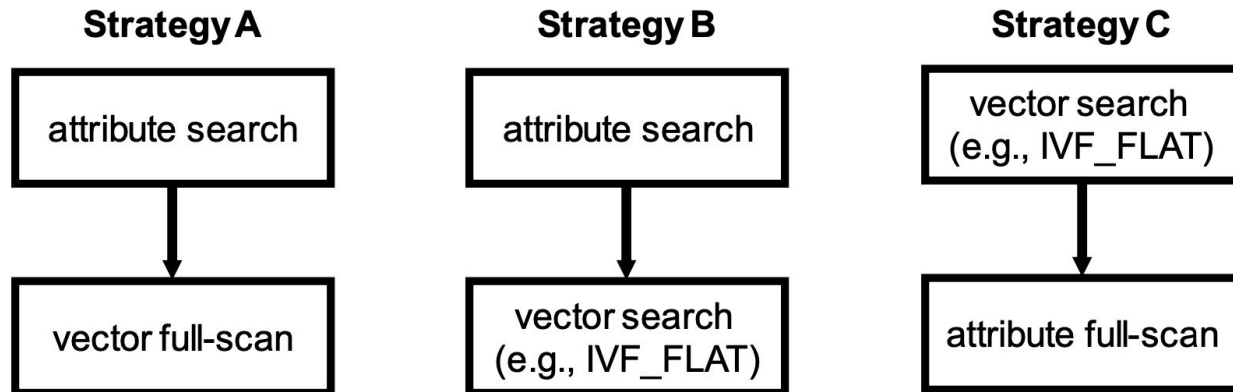
- New SortPlan in algebra tree
 - Needs to estimate its own cost

Cost Estimation (IVF_FLAT)

- To select top clusters: $B(\text{centroid file})$
- Scan for each cluster: $B(\text{data file of a centroid})$

Attribute Filtering

```
SELECT p.id,  
       p.text_vector <-> '...' AS dist  
FROM posts AS p  
WHERE p.date < '...'  
ORDER BY dist ASC LIMIT 10;
```



- Best strategy determined by estimated costs

Outline


- Why Vector DBMS?
- AKNN Search Algorithms
- Challenges at System-level
- Case study: PASE (System R-like)
 - Data model & Query Format
 - Index Building & Update
 - Planning & Cost Estimation
- **Case study: Milvus (purpose-built)**
 - Storage & Consistency Model
 - Computing & Threads
 - Query Algorithms

Milvus

- “Milvus: A Purpose-Built Vector Data Management System,” in SIGMOD’21
 - A dedicated system
- Pros:
 - Supports disk storage, dynamic data, attribute filtering
 - Much higher performance than PASE

Storage

- Column storage based on Log Structured Merge (LSM) tree

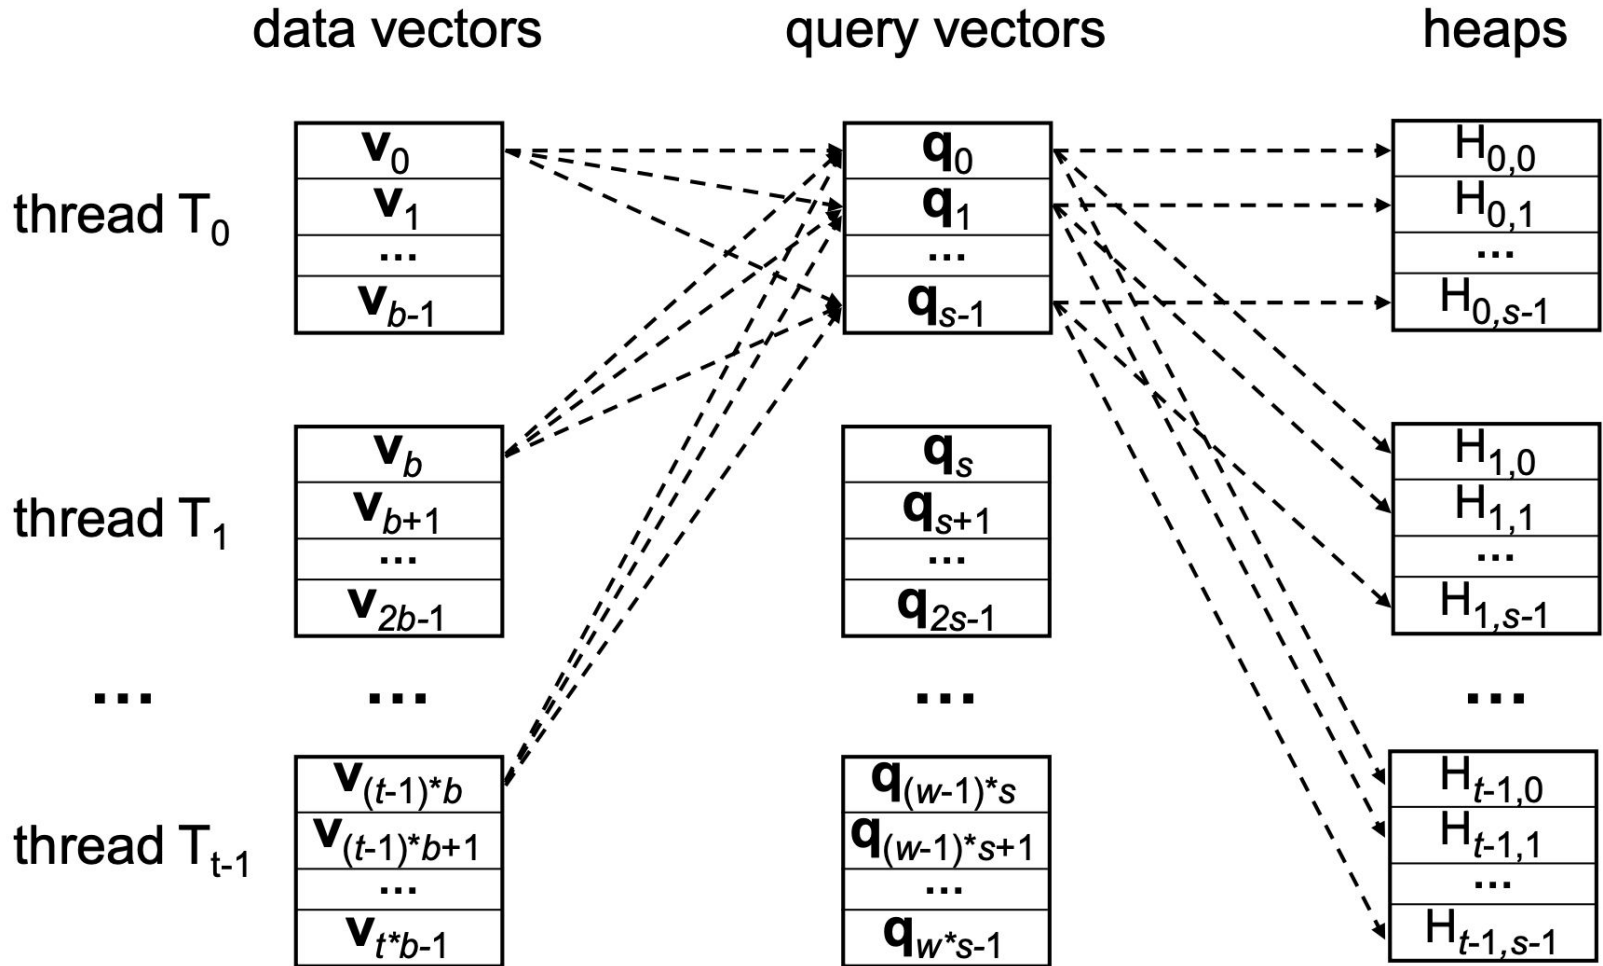


- Out-of-place updates
- SSTables (segments) are the unit of buffering/searching

Consistency Model

- Snapshot isolation

Thread Model



Query Planning

